

# Orleans

Microsoft's Distributed Systems framework for .NET

Reuben Bond



# Agenda

**01**

The hidden cost of statelessness

**02**

The stateful solution

**03**

Microsoft Orleans

**04**

What's new since Orleans 8.0

**05**

Placement & load balancing

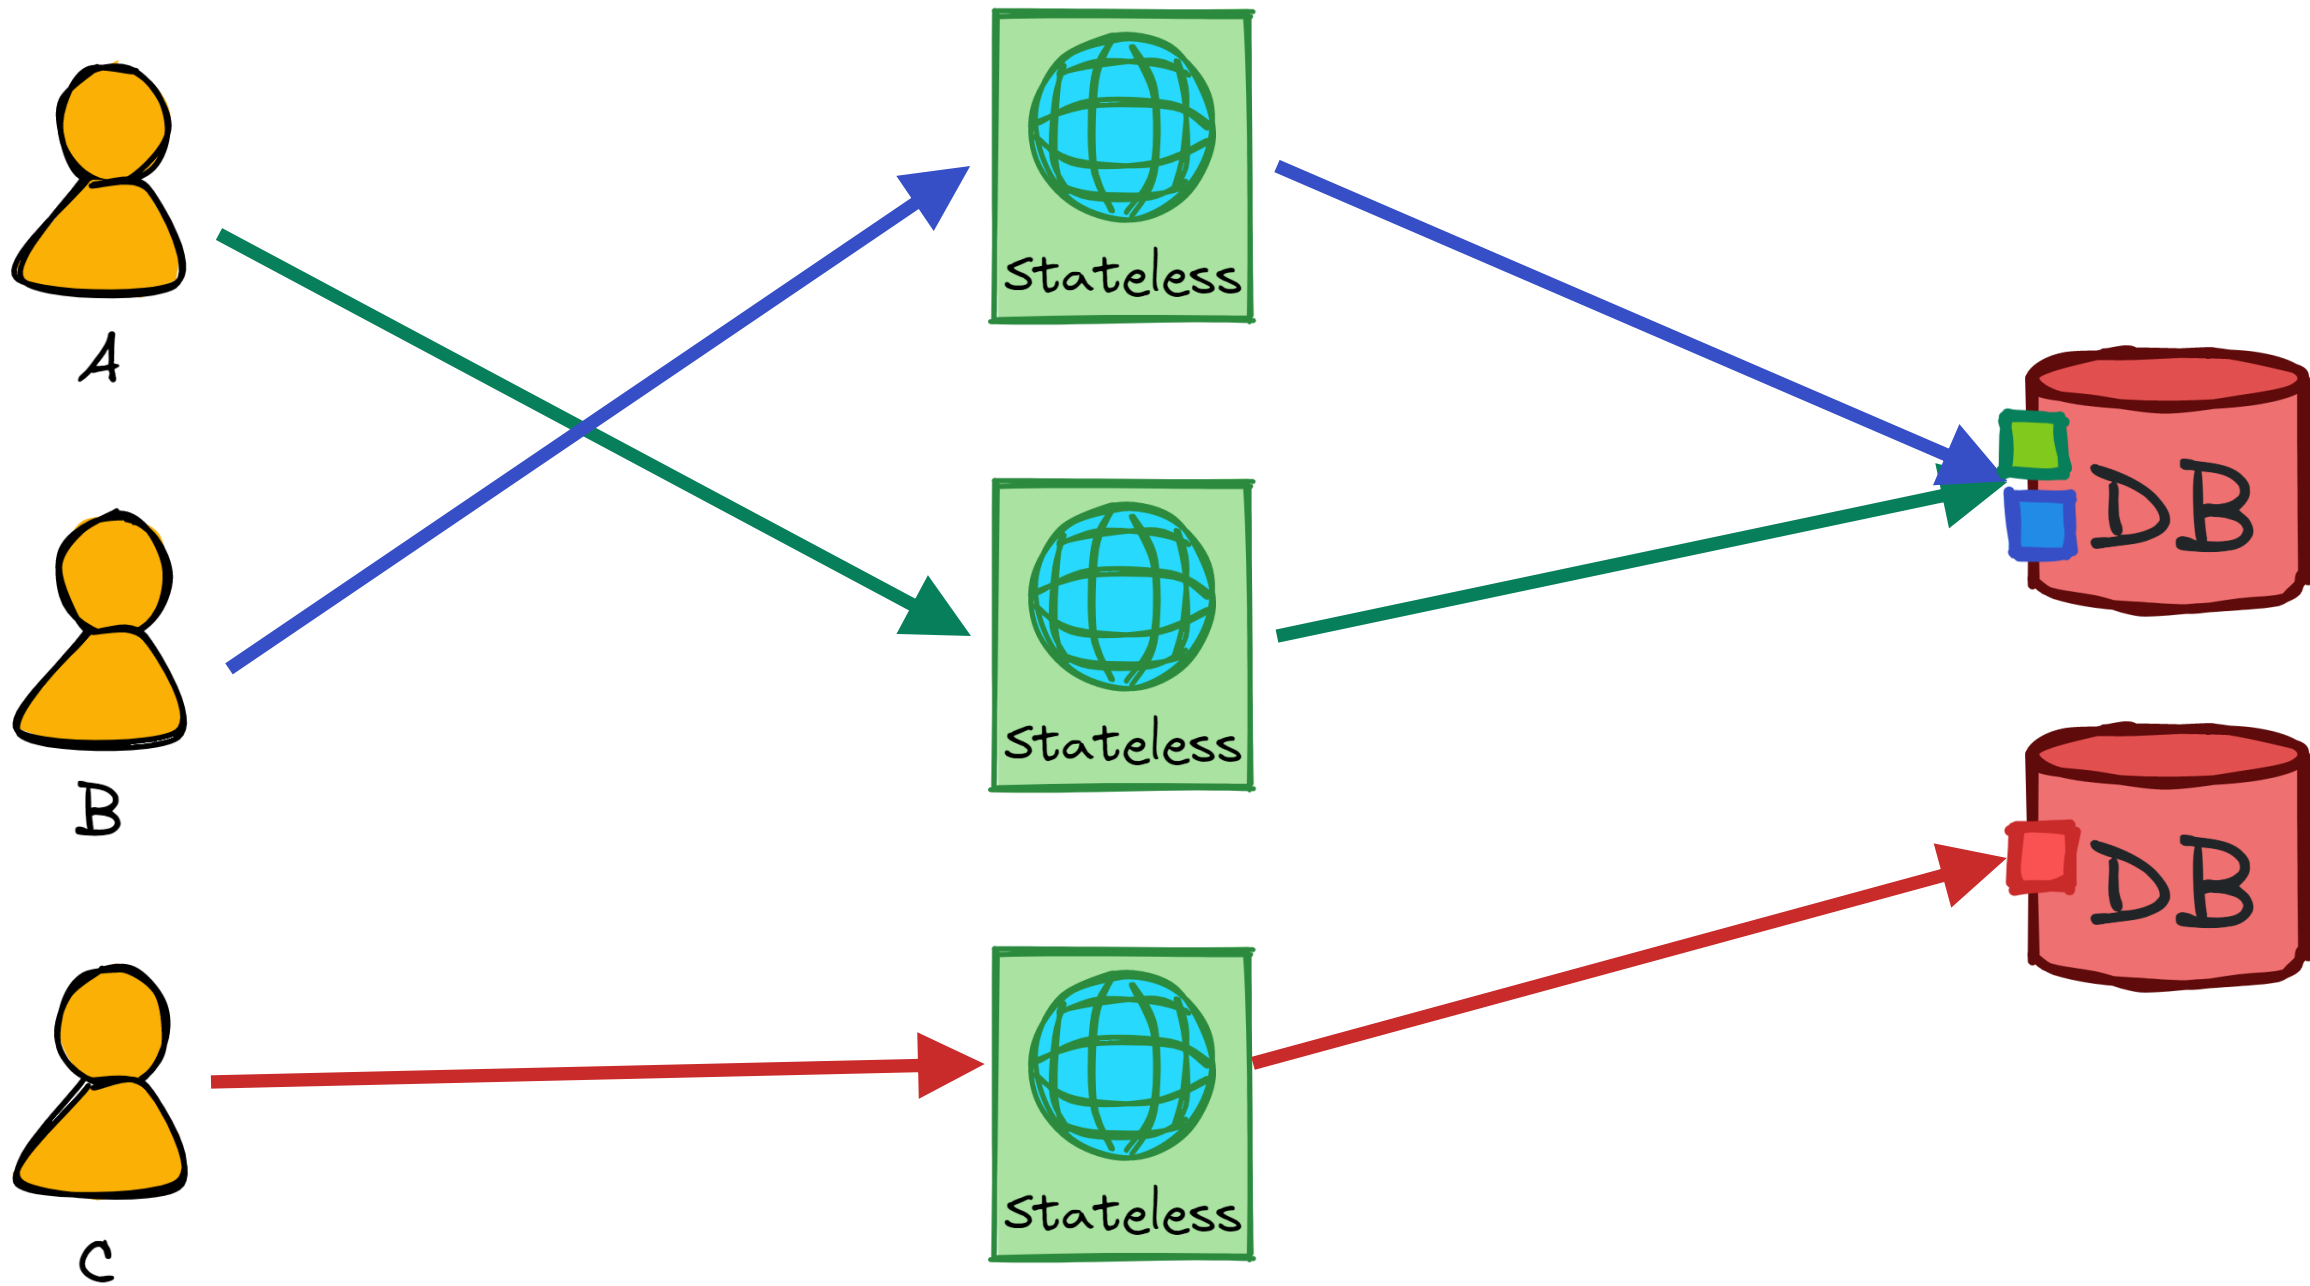
**06**

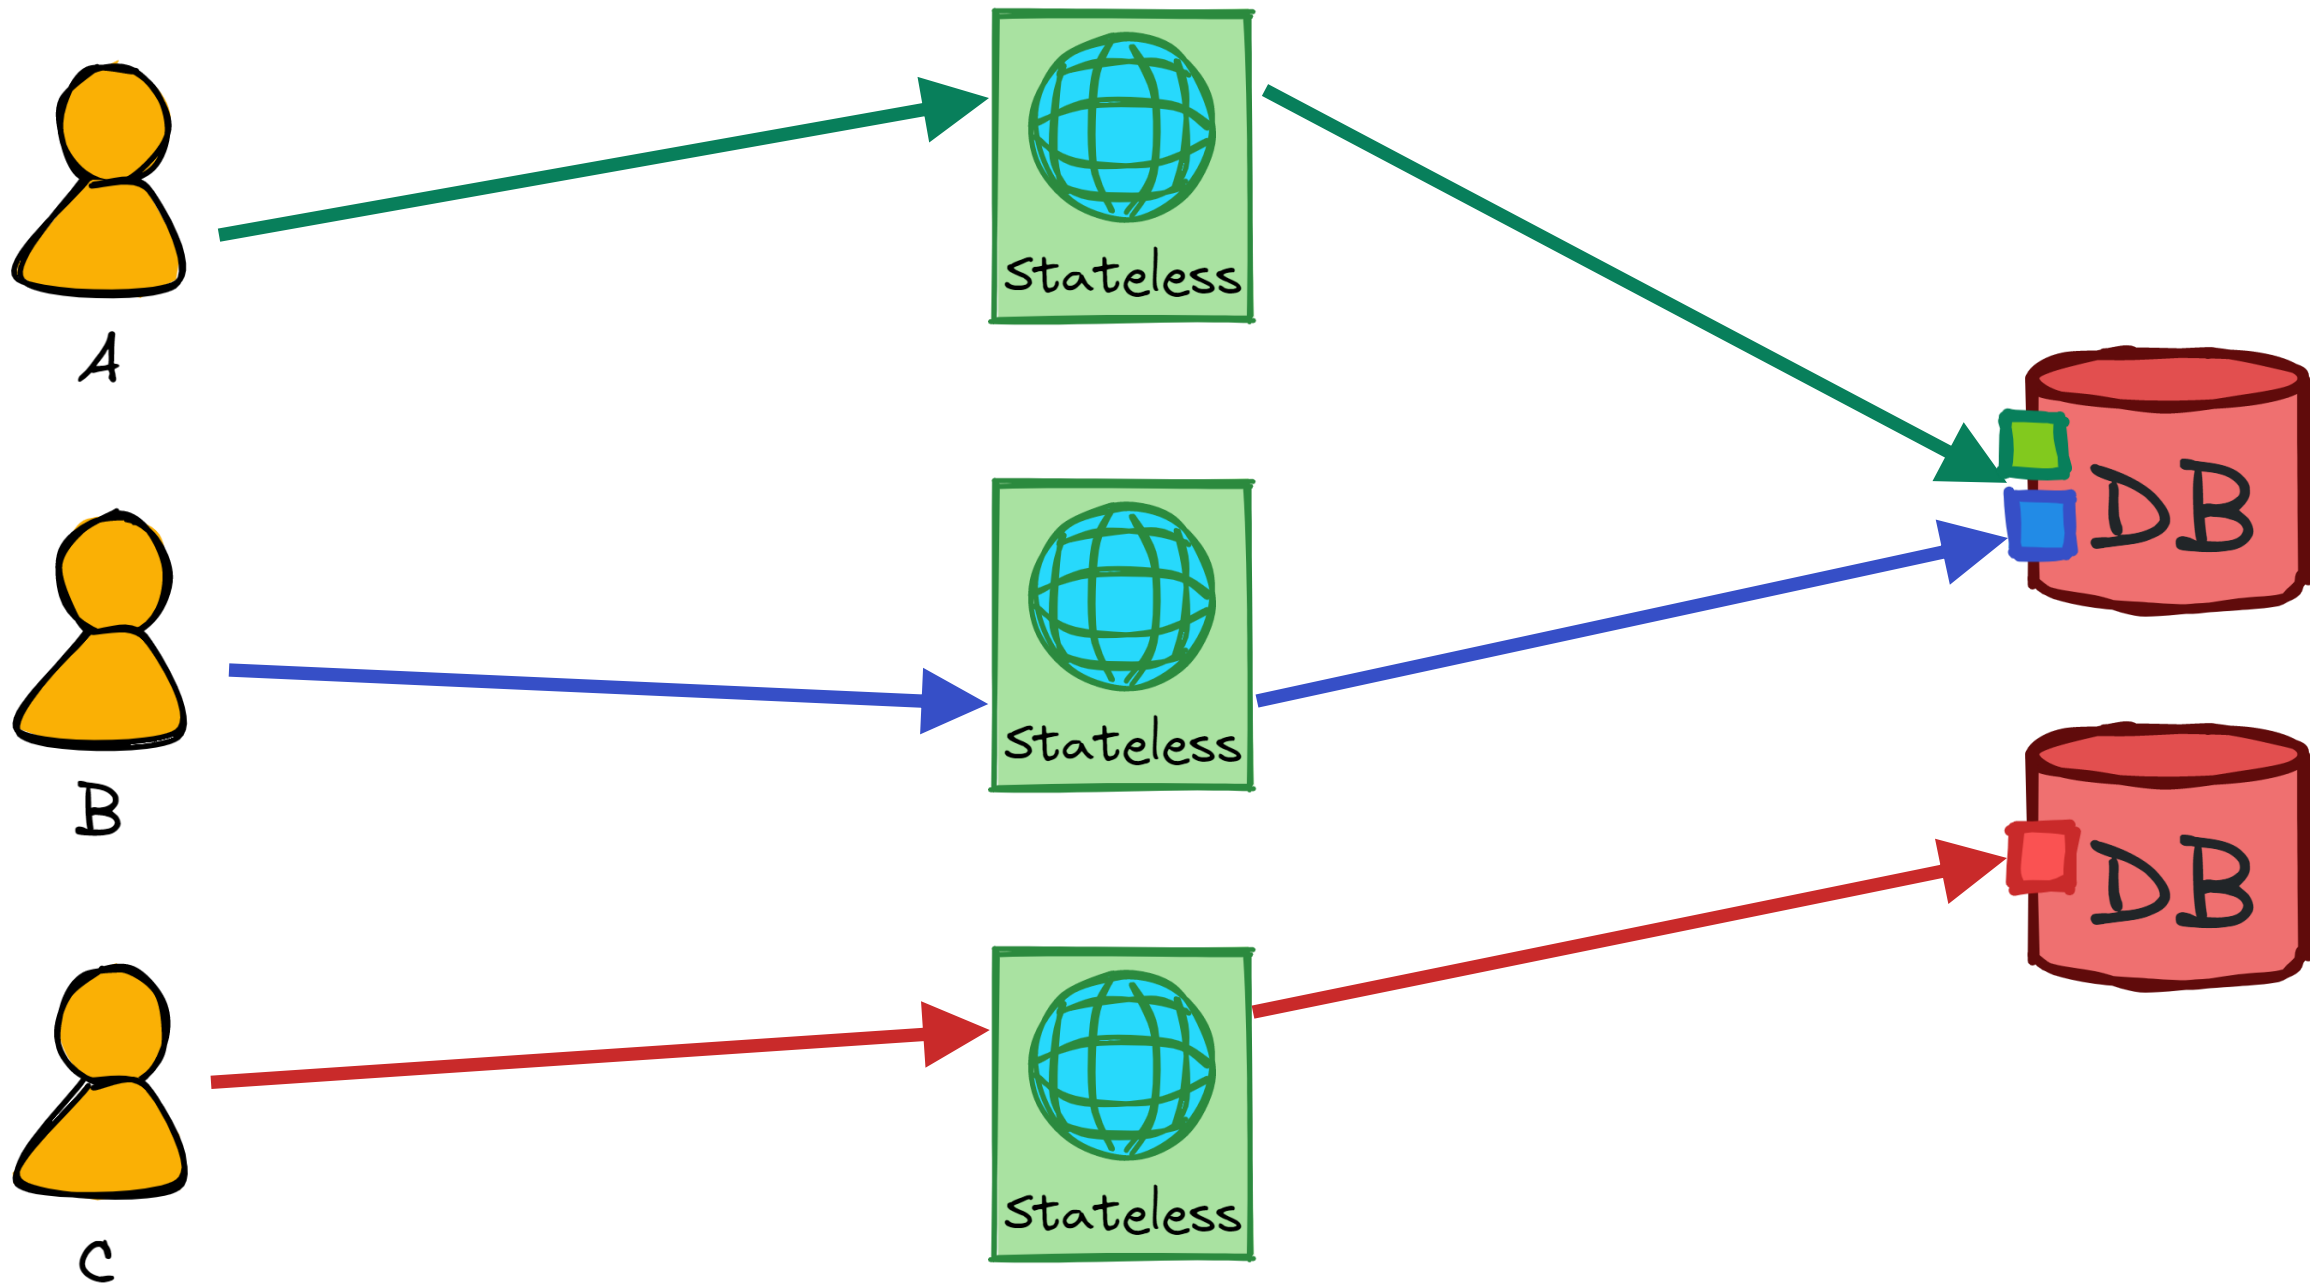
Q&A

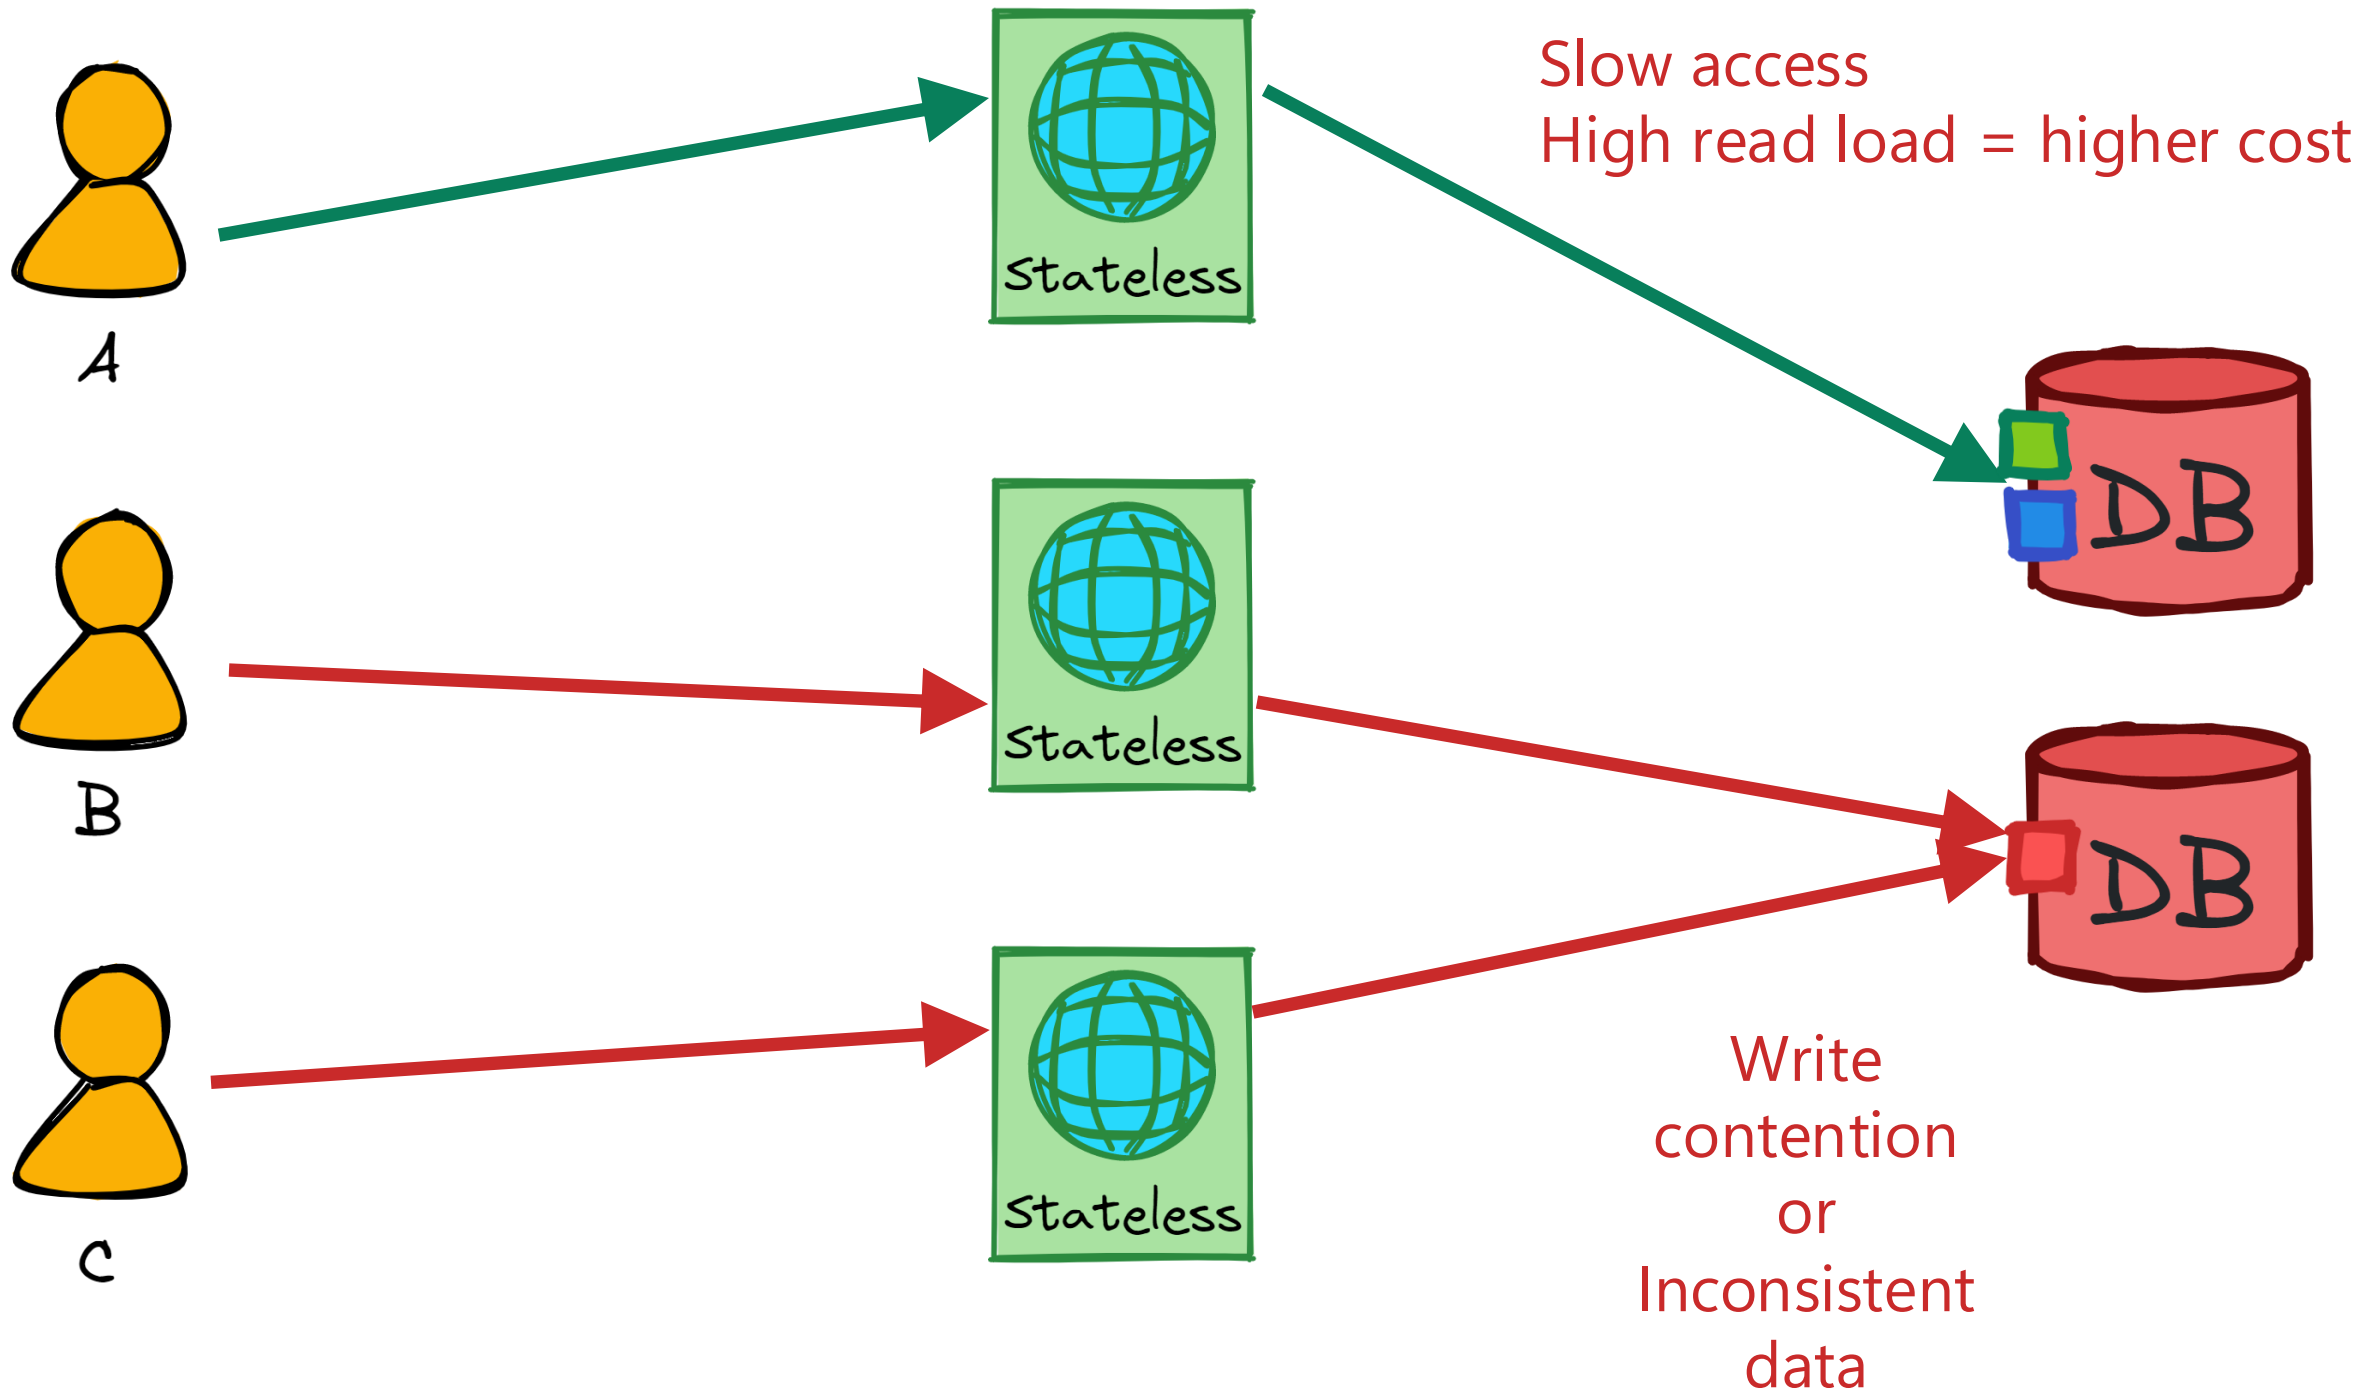
Abstract purple geometric shapes including a large curved band at the top right, a solid sphere in the center, and a thin ring at the bottom left.

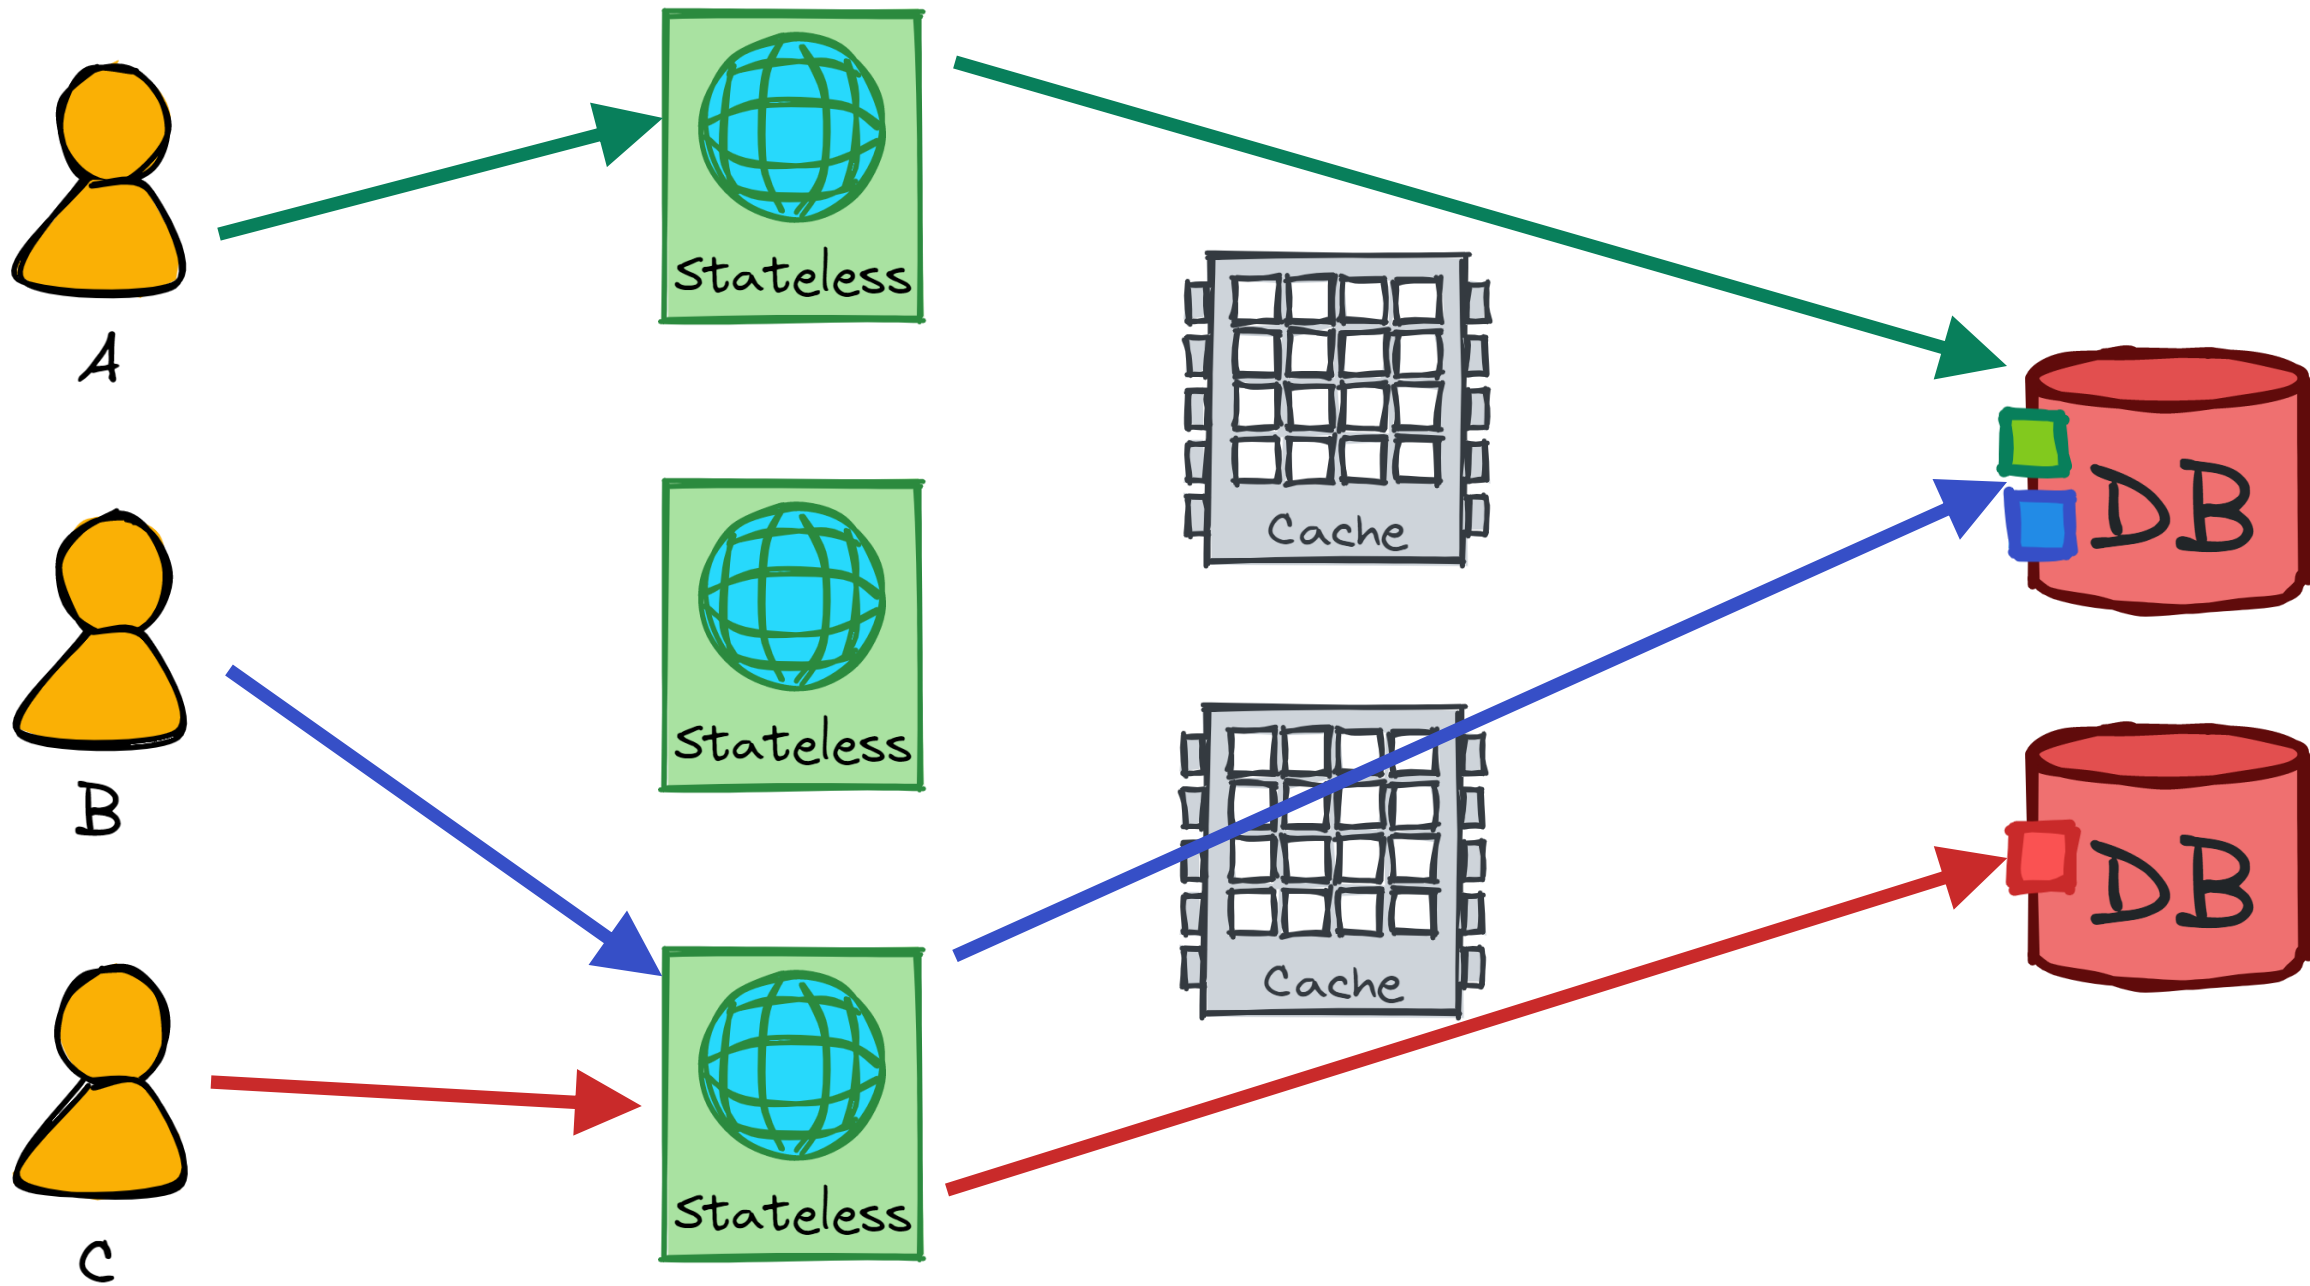
01

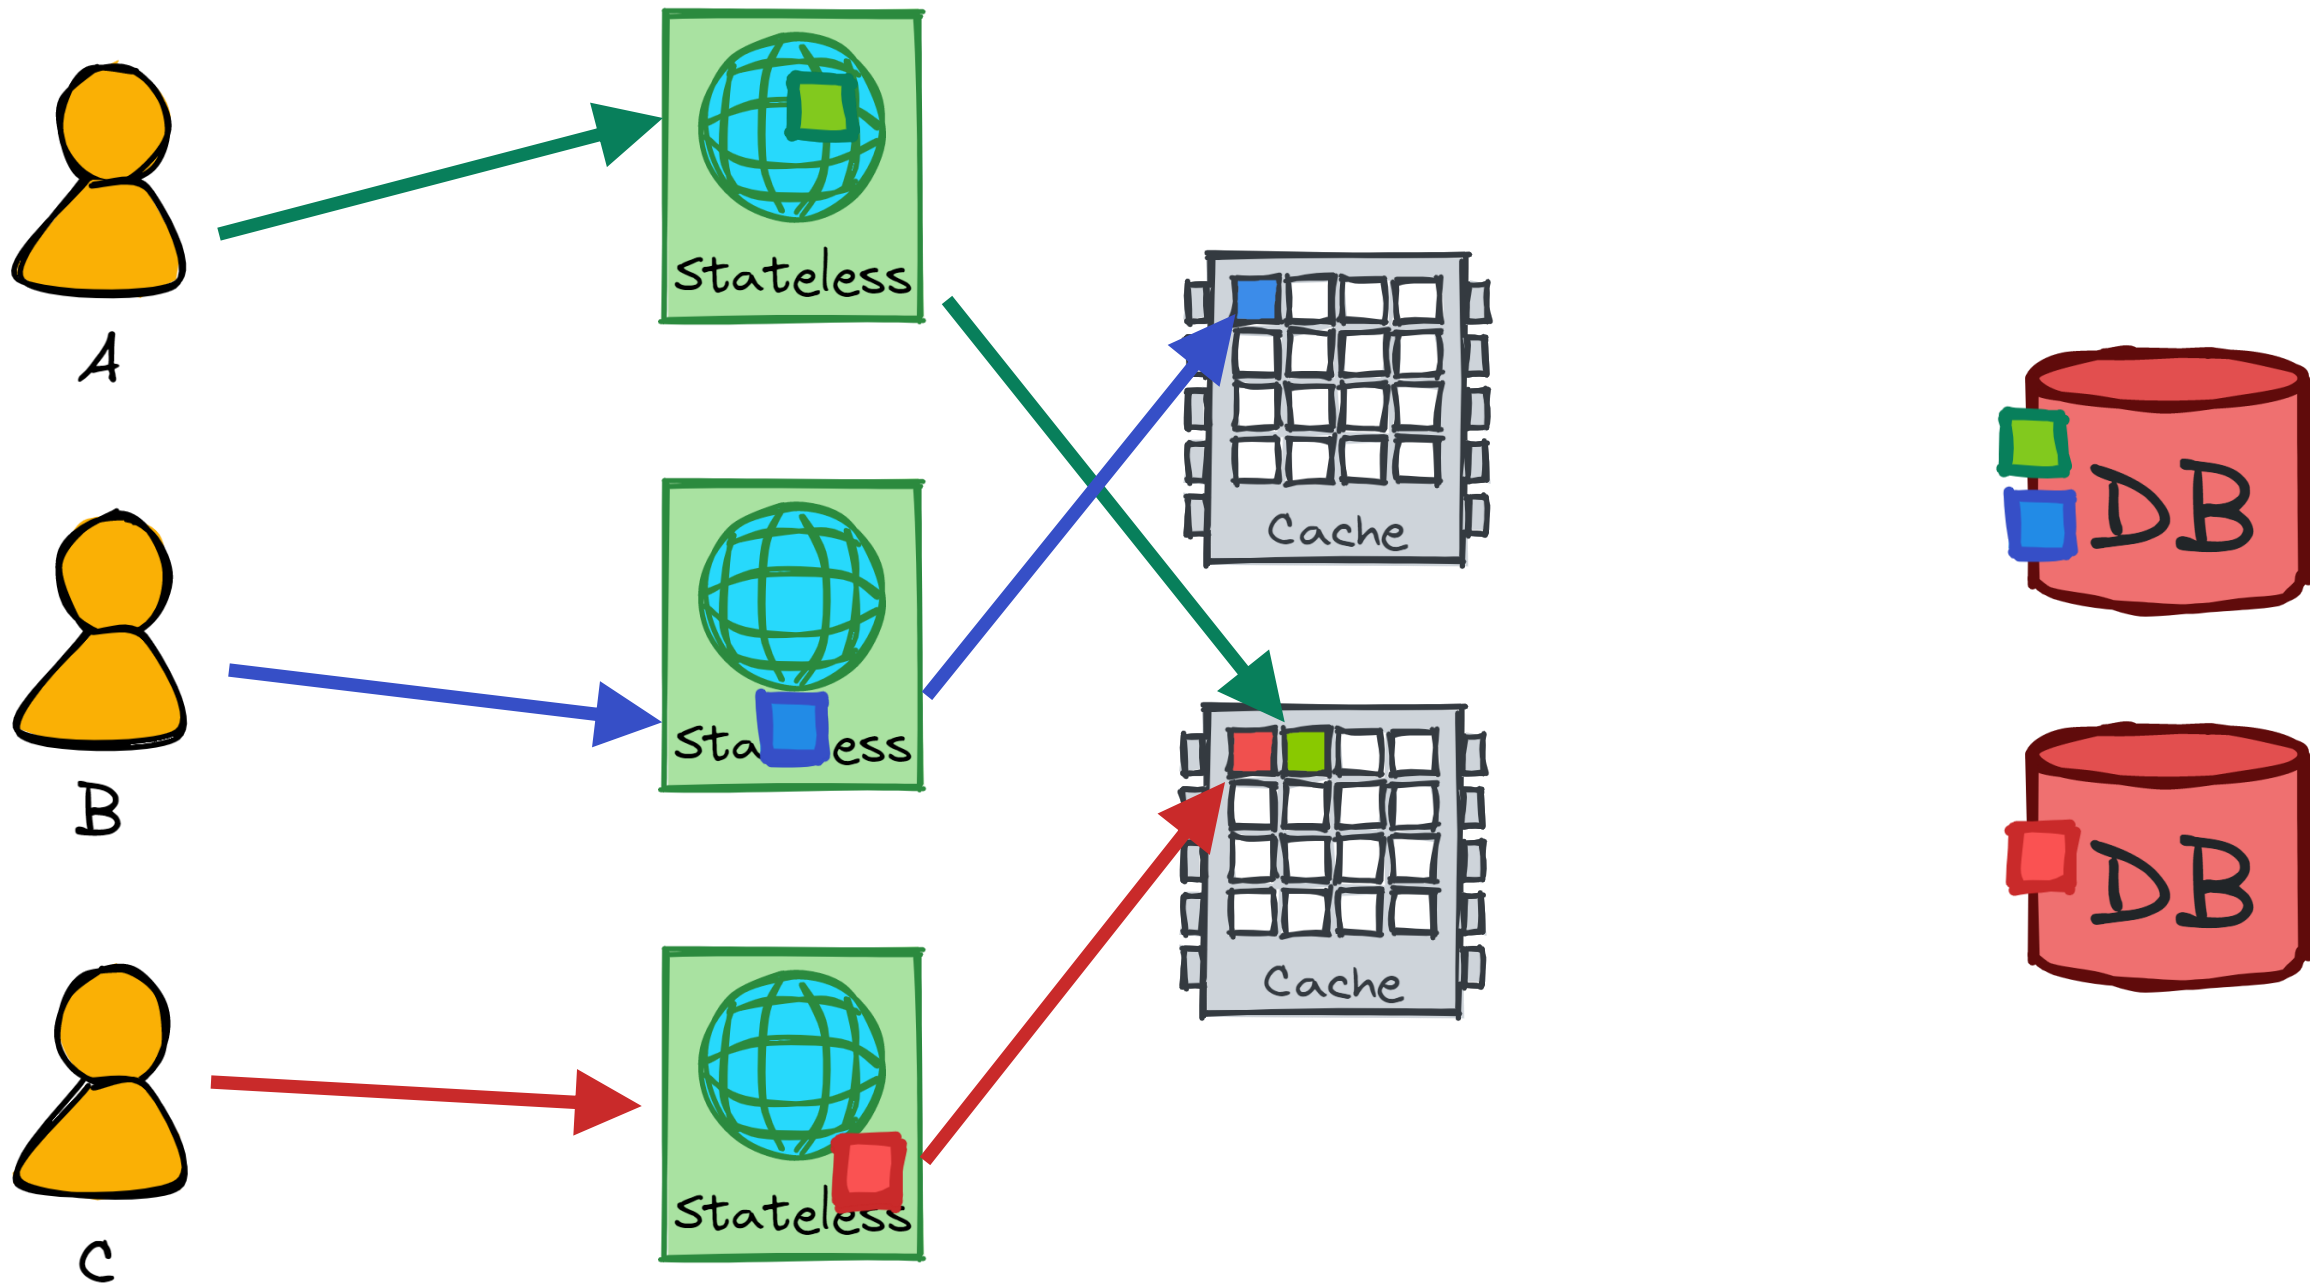
## **The hidden cost of statelessness**



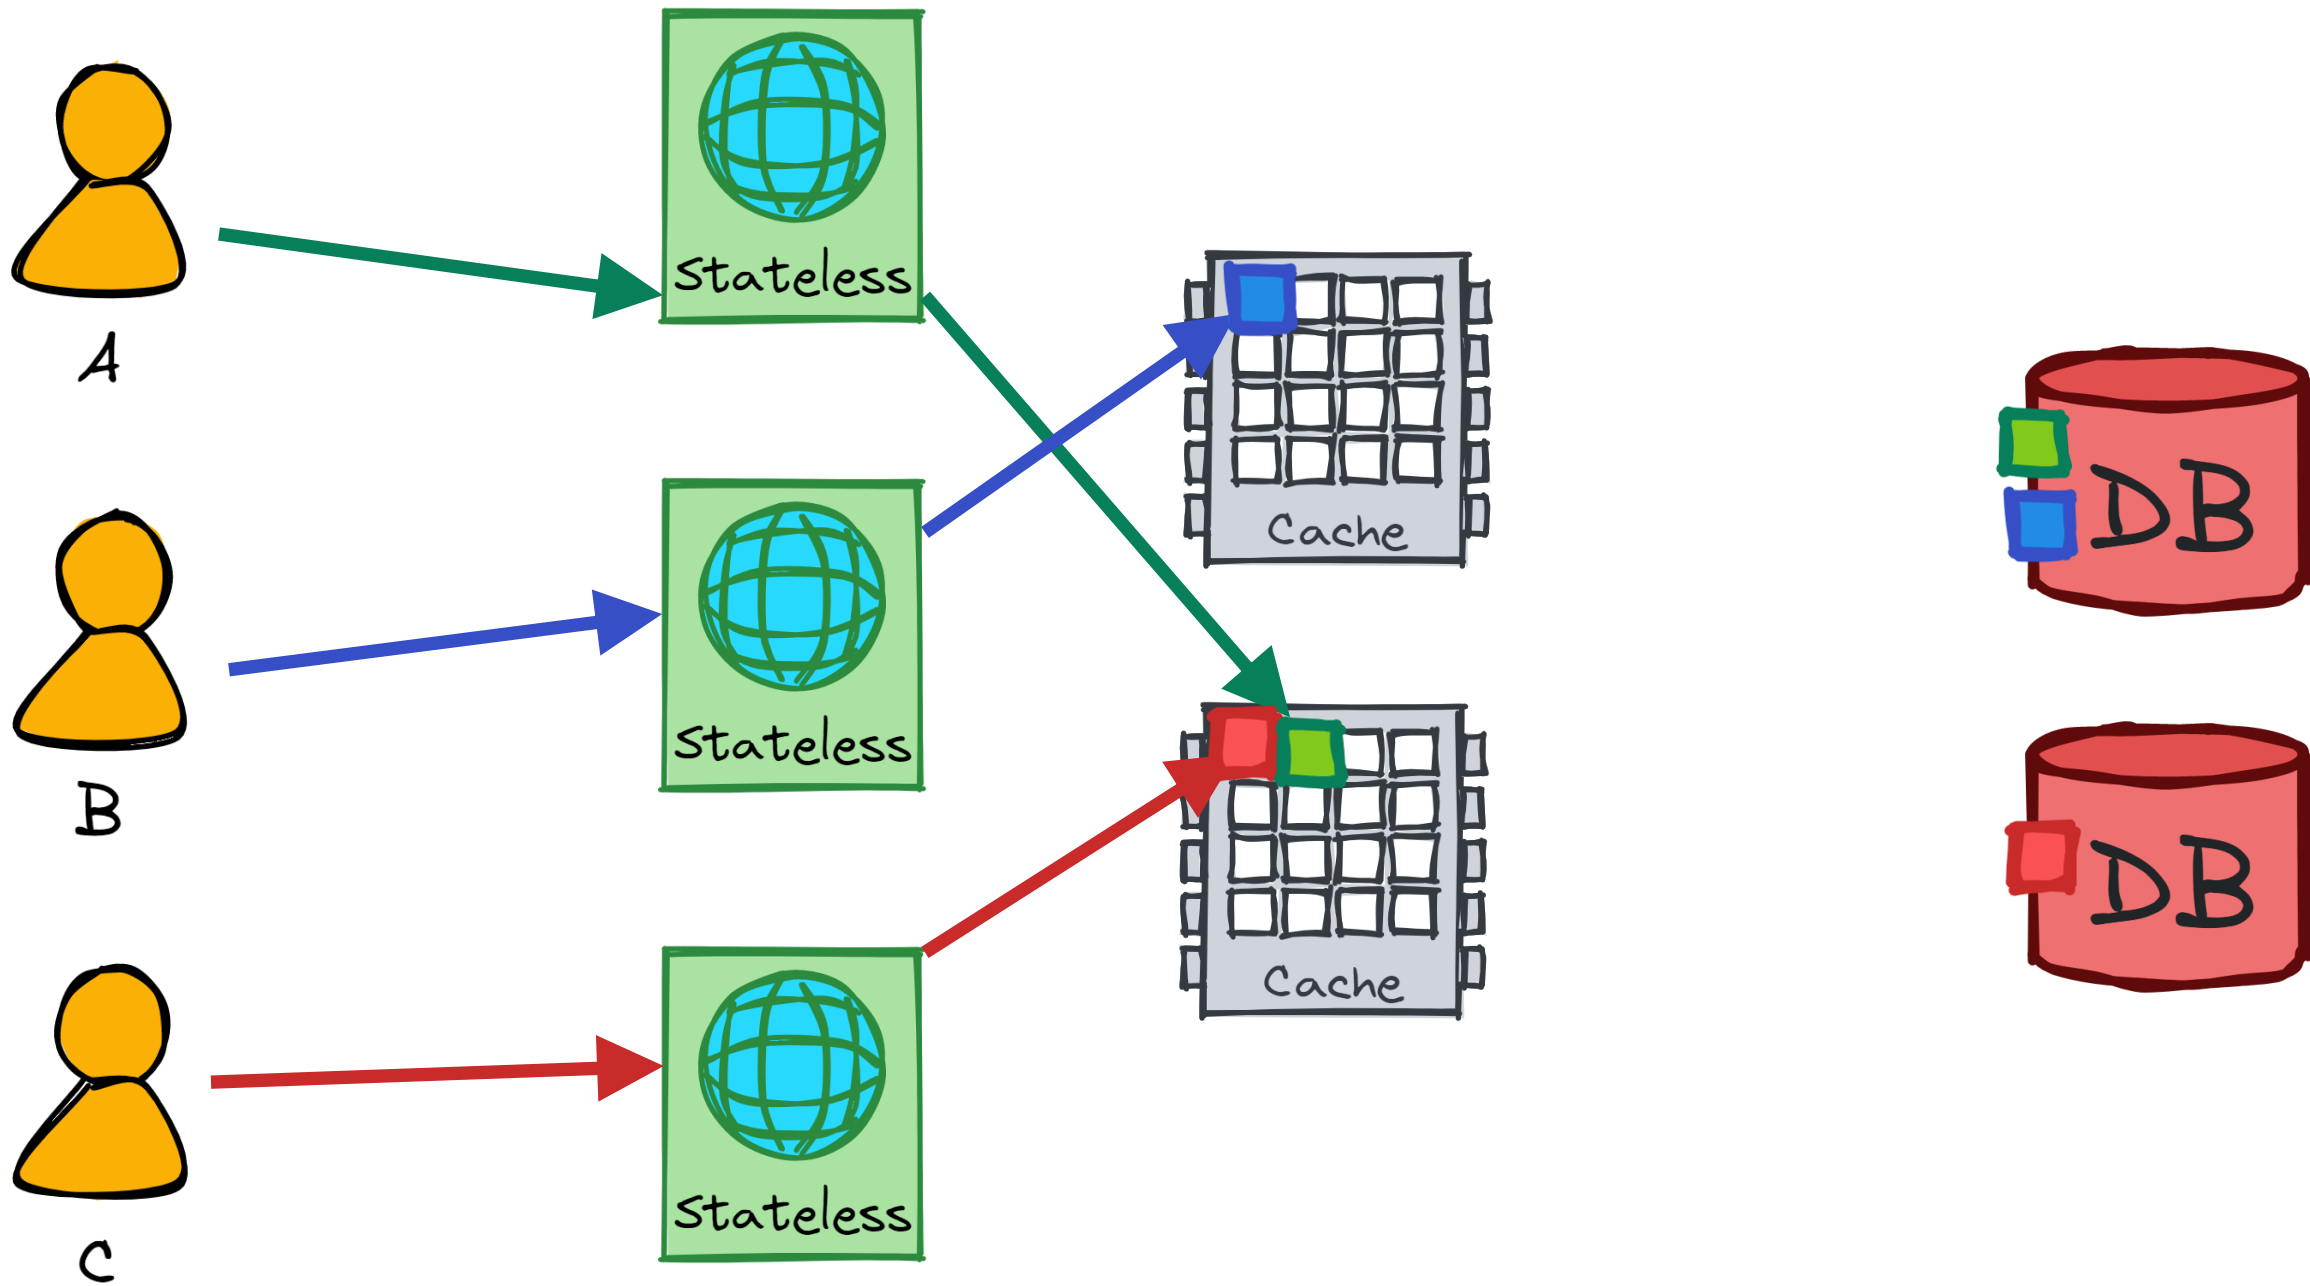


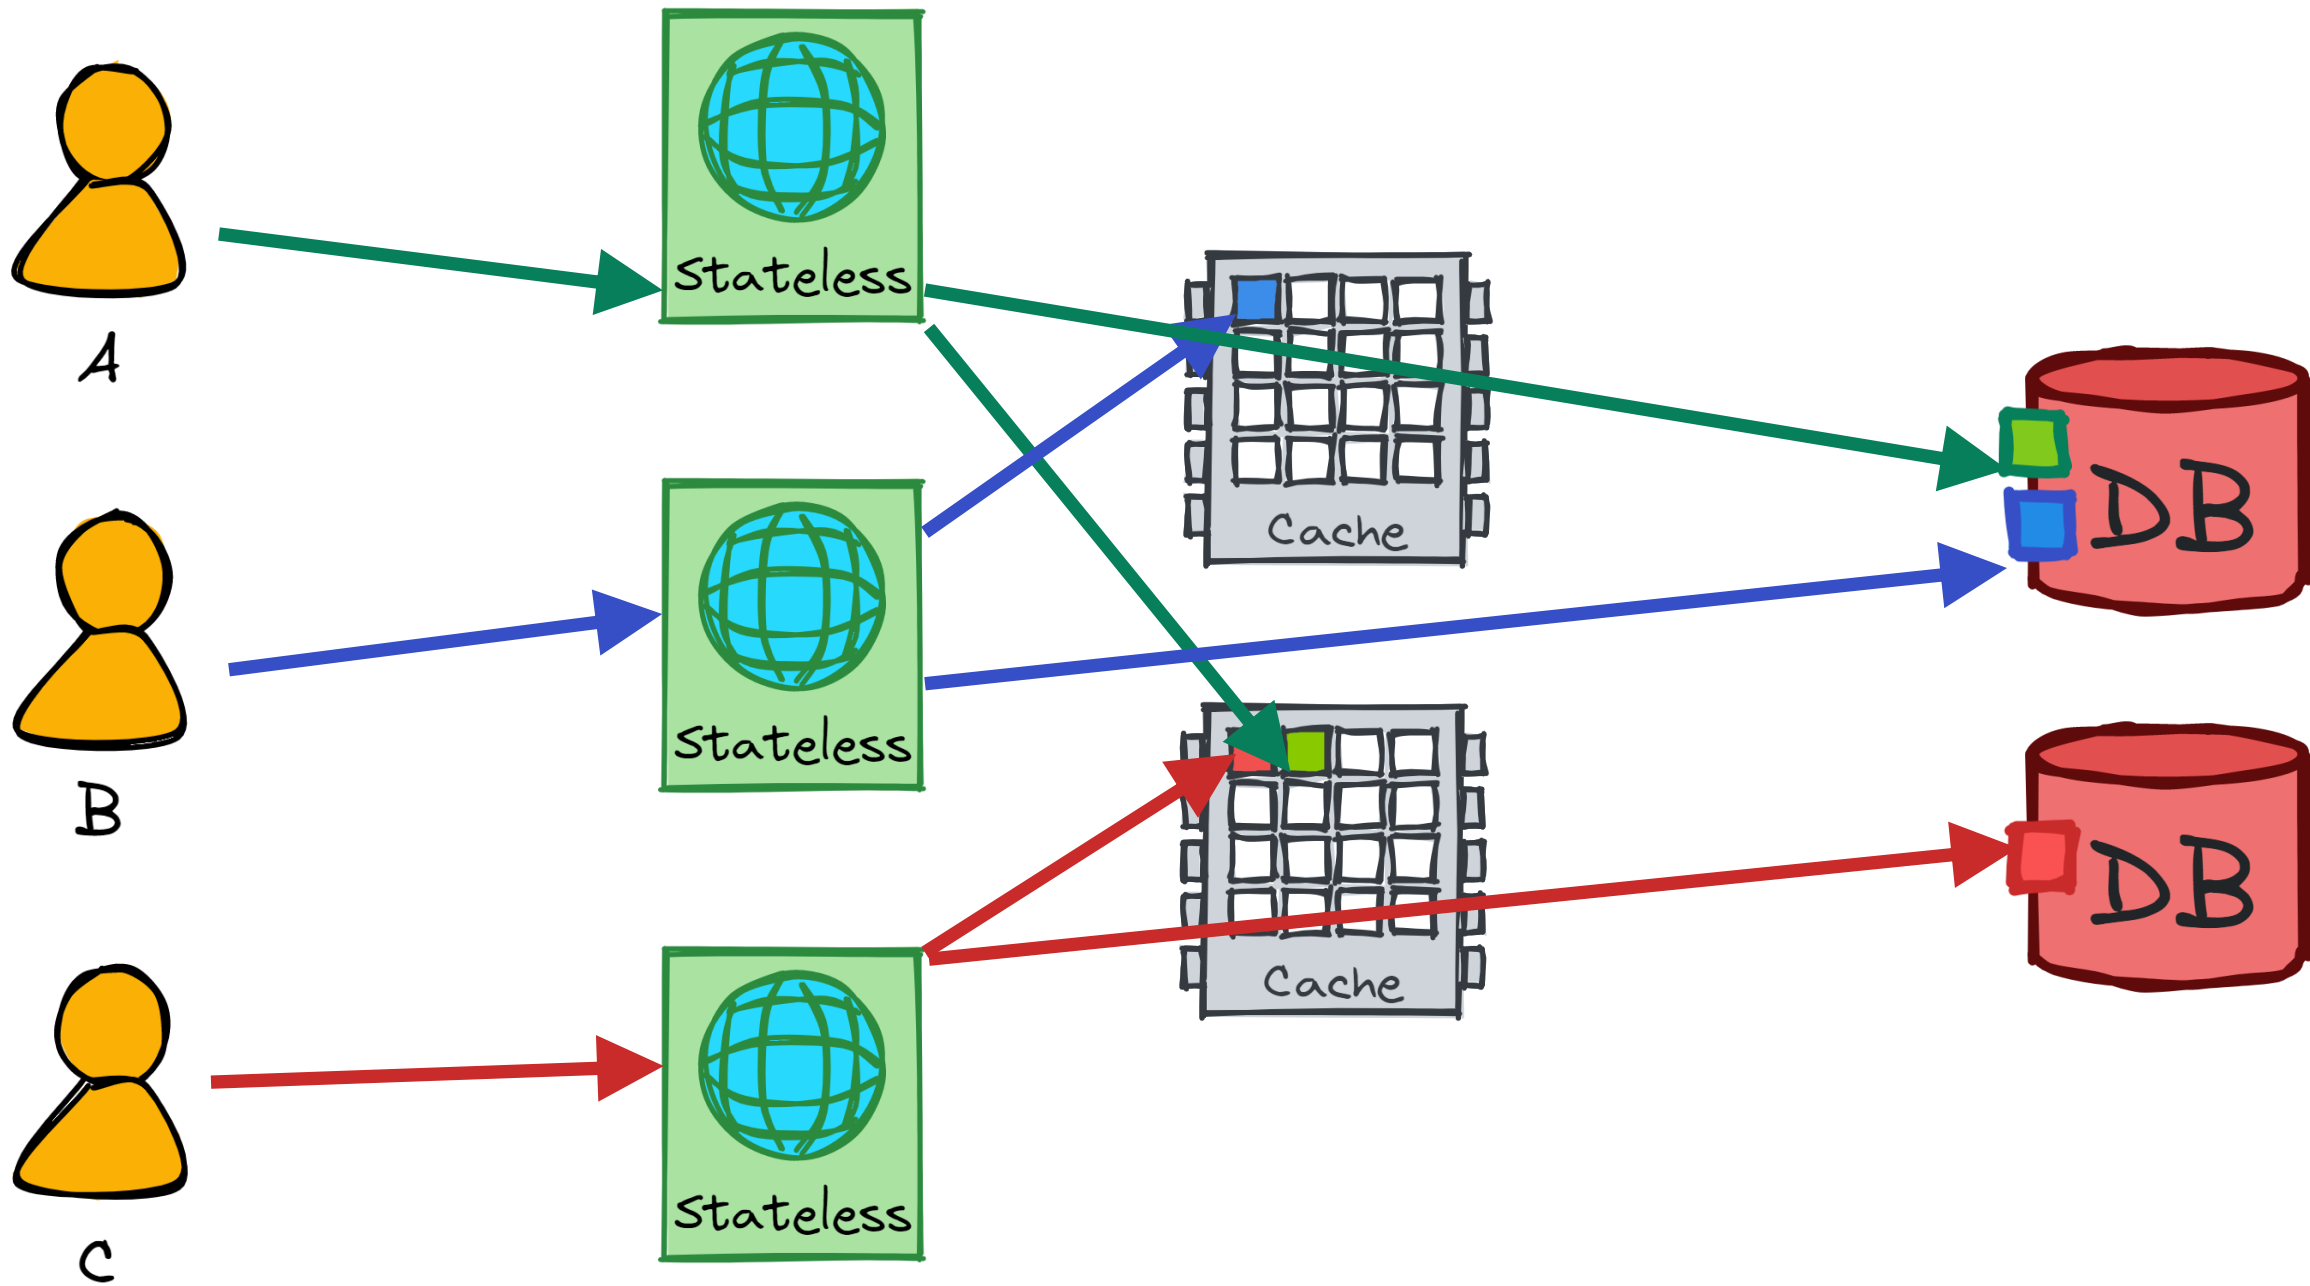


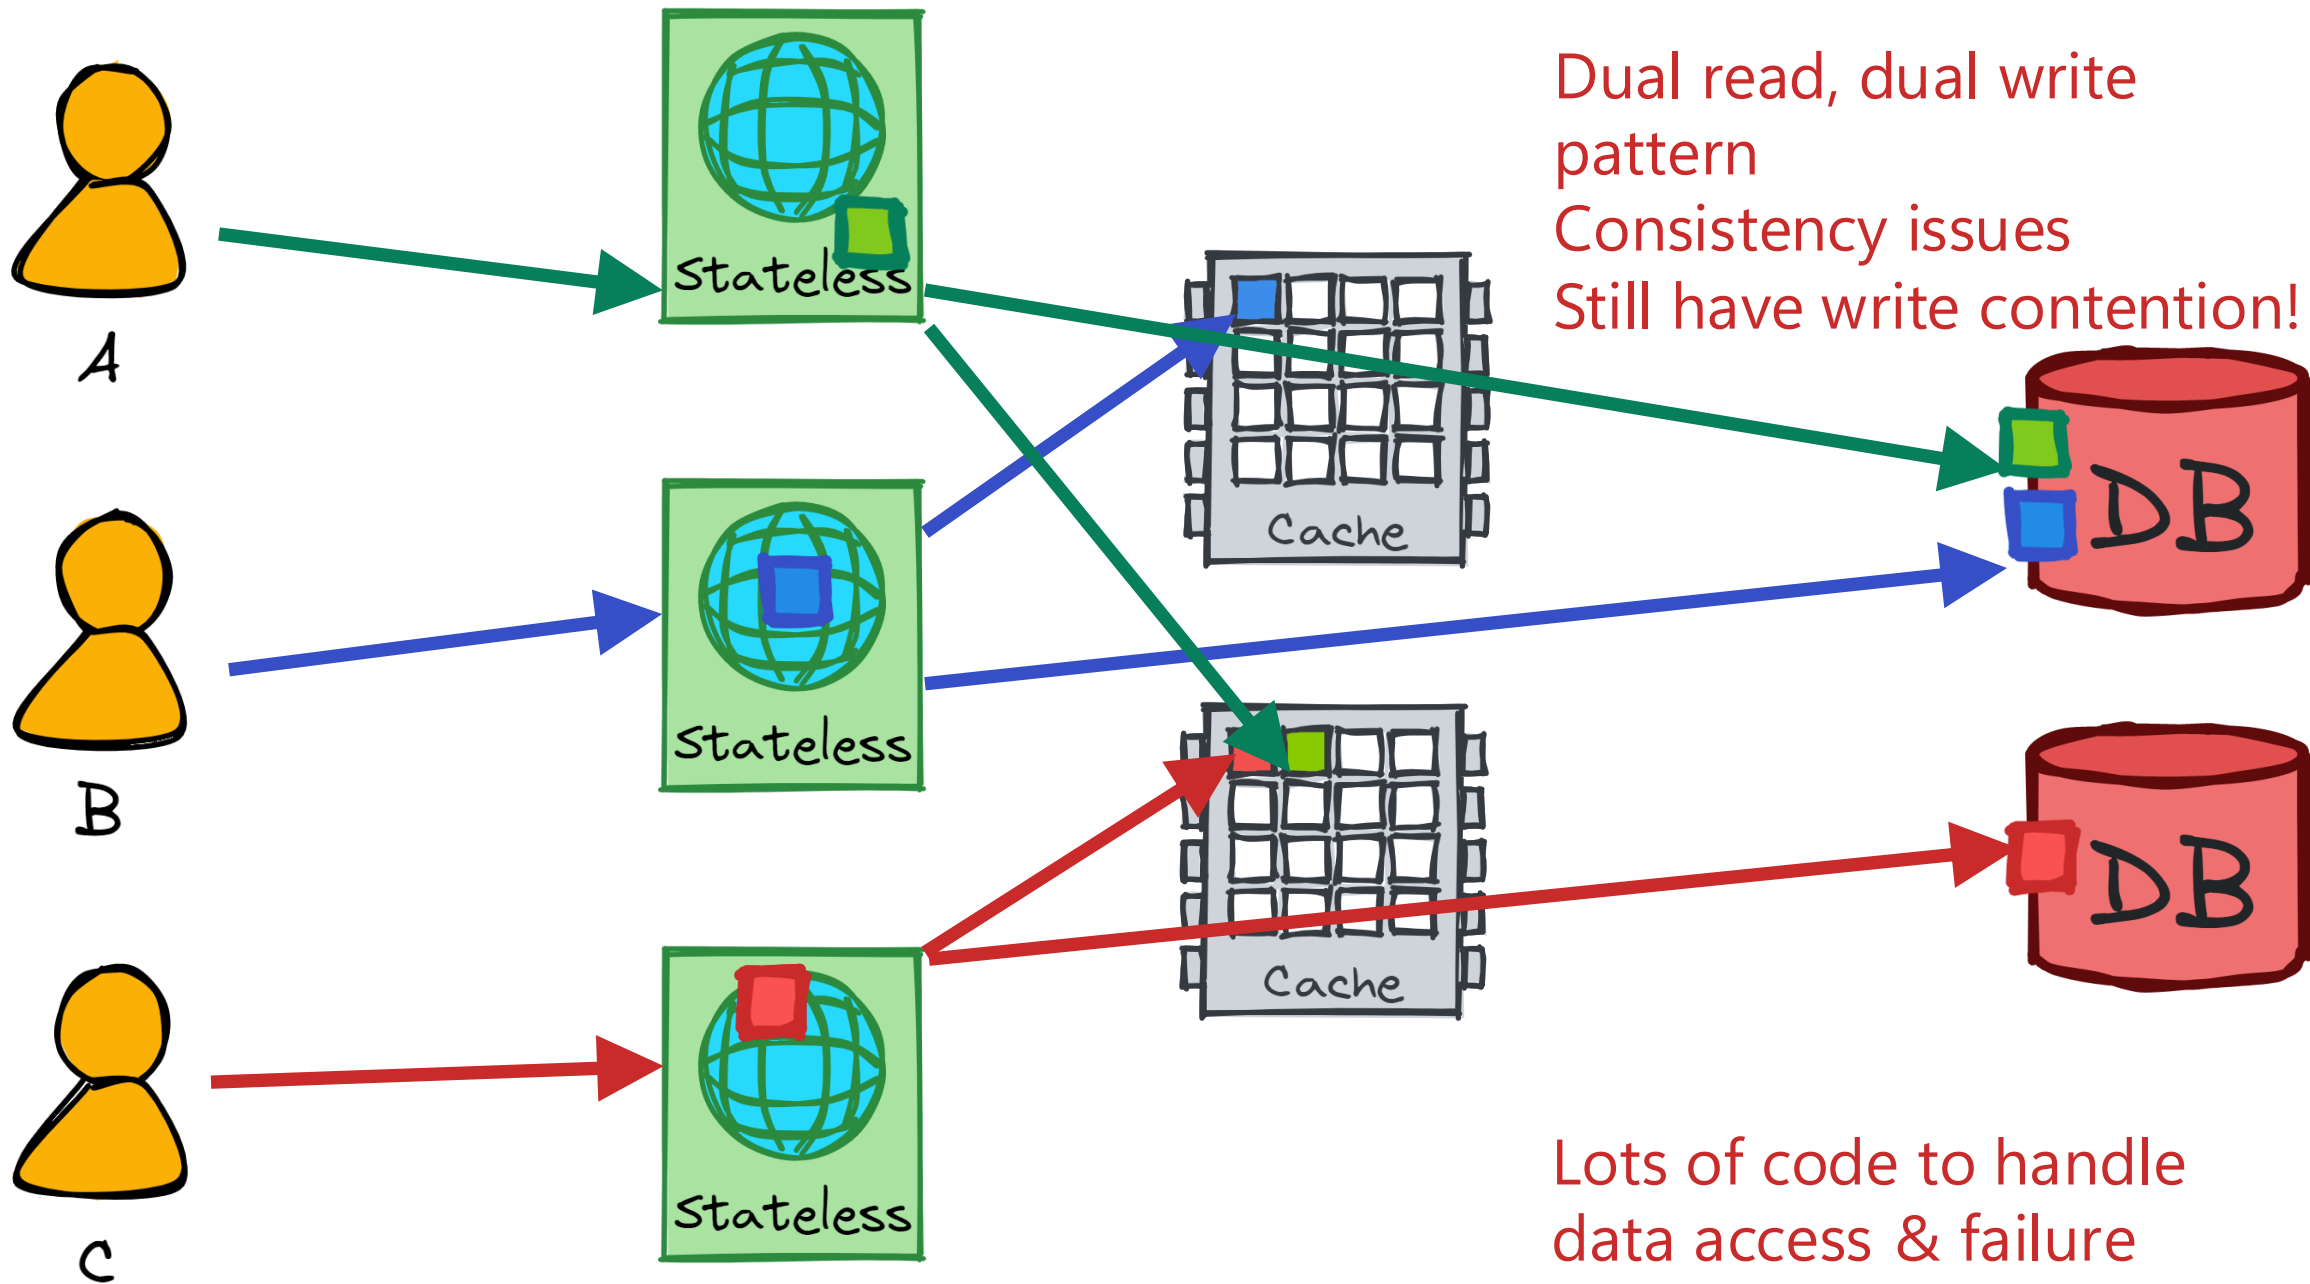












Lots of code to handle  
data access & failure  
cases!

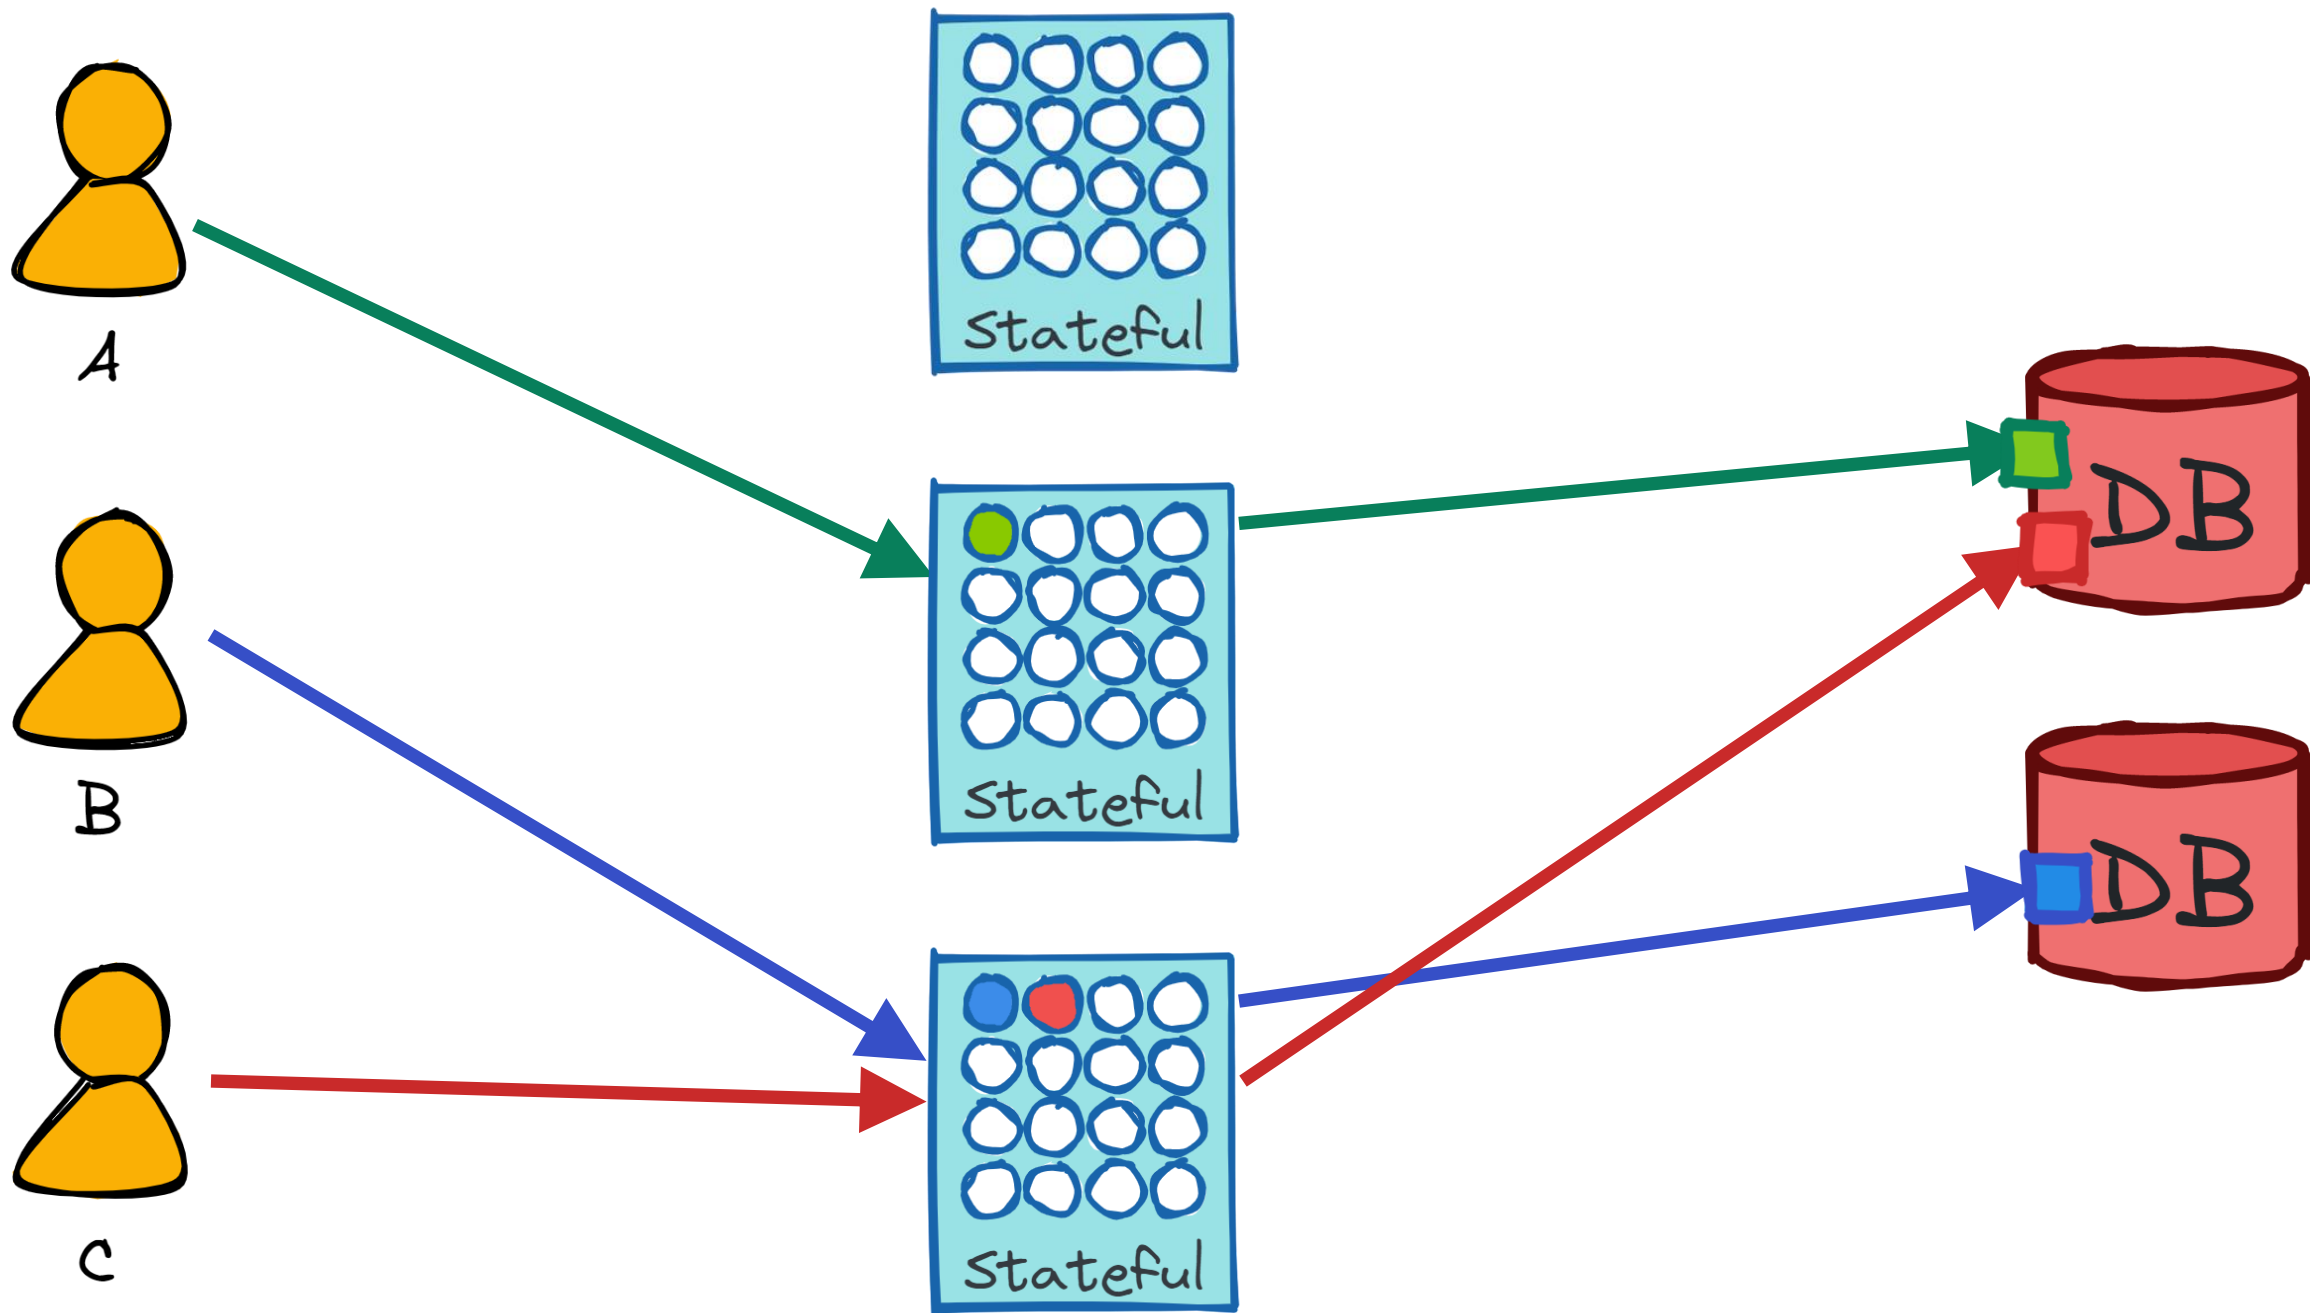
# The hidden costs of statelessness

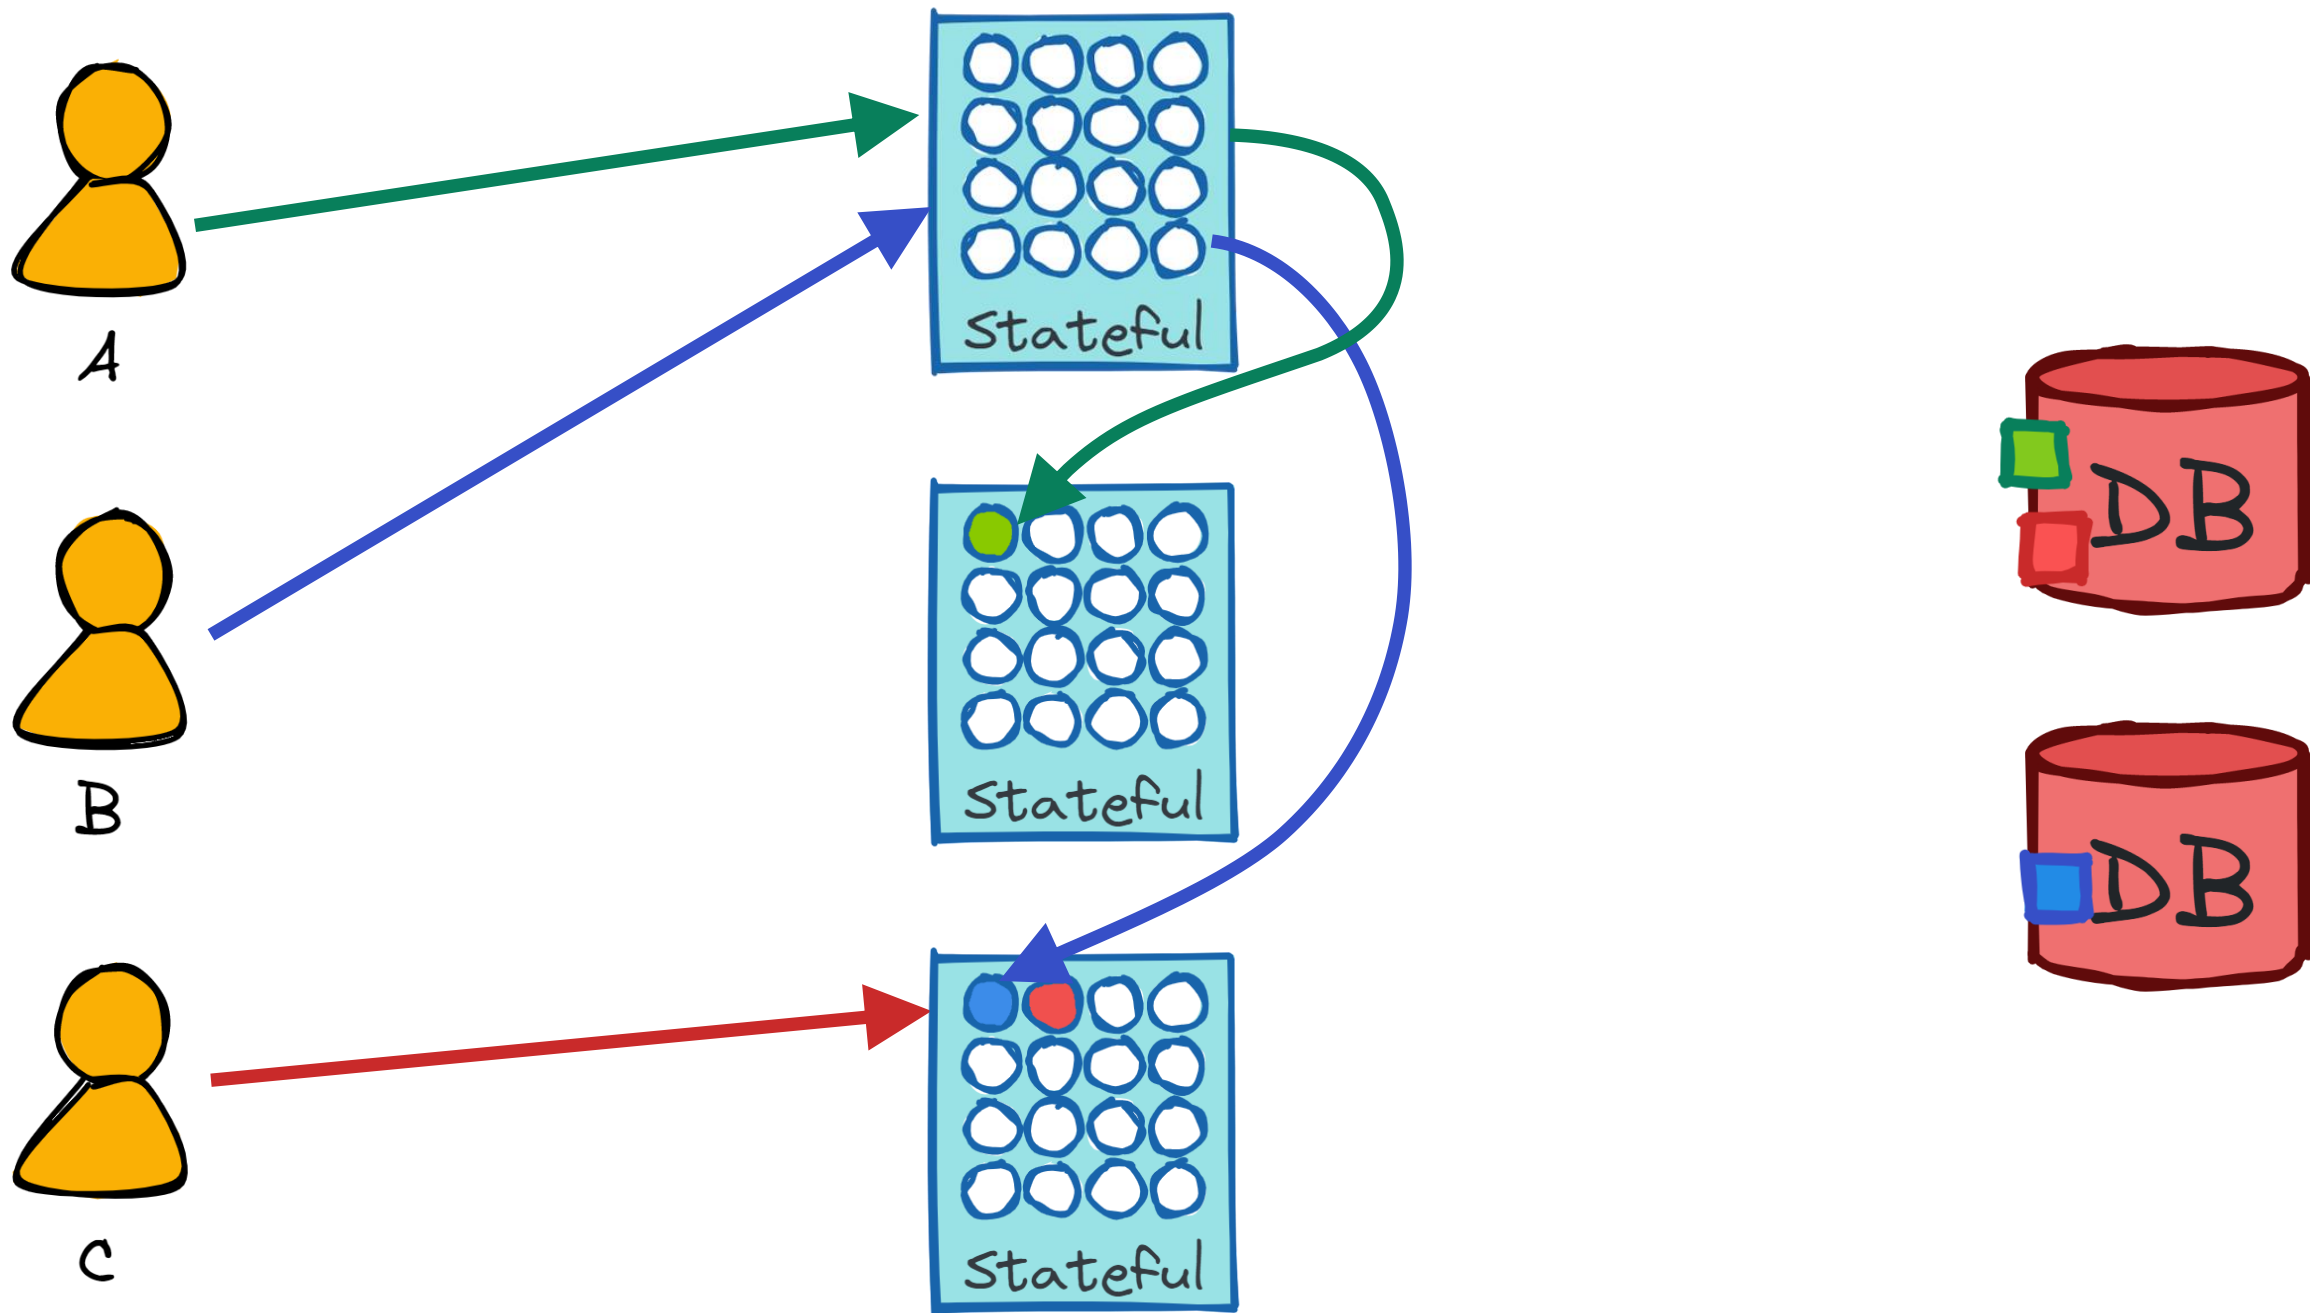
- Higher latency
  - Read, modify, write = 2+ I/O to cache/DB per request
  - Contention retries/backoff
- Lower throughput
  - (De)serialization cost for cache, I/O costs, read amplification, wasted I/O to resolve write contention
- Higher database cost, poor scalability
- Cache infrastructure costs
- Cache coherence complexities
- More code to maintain

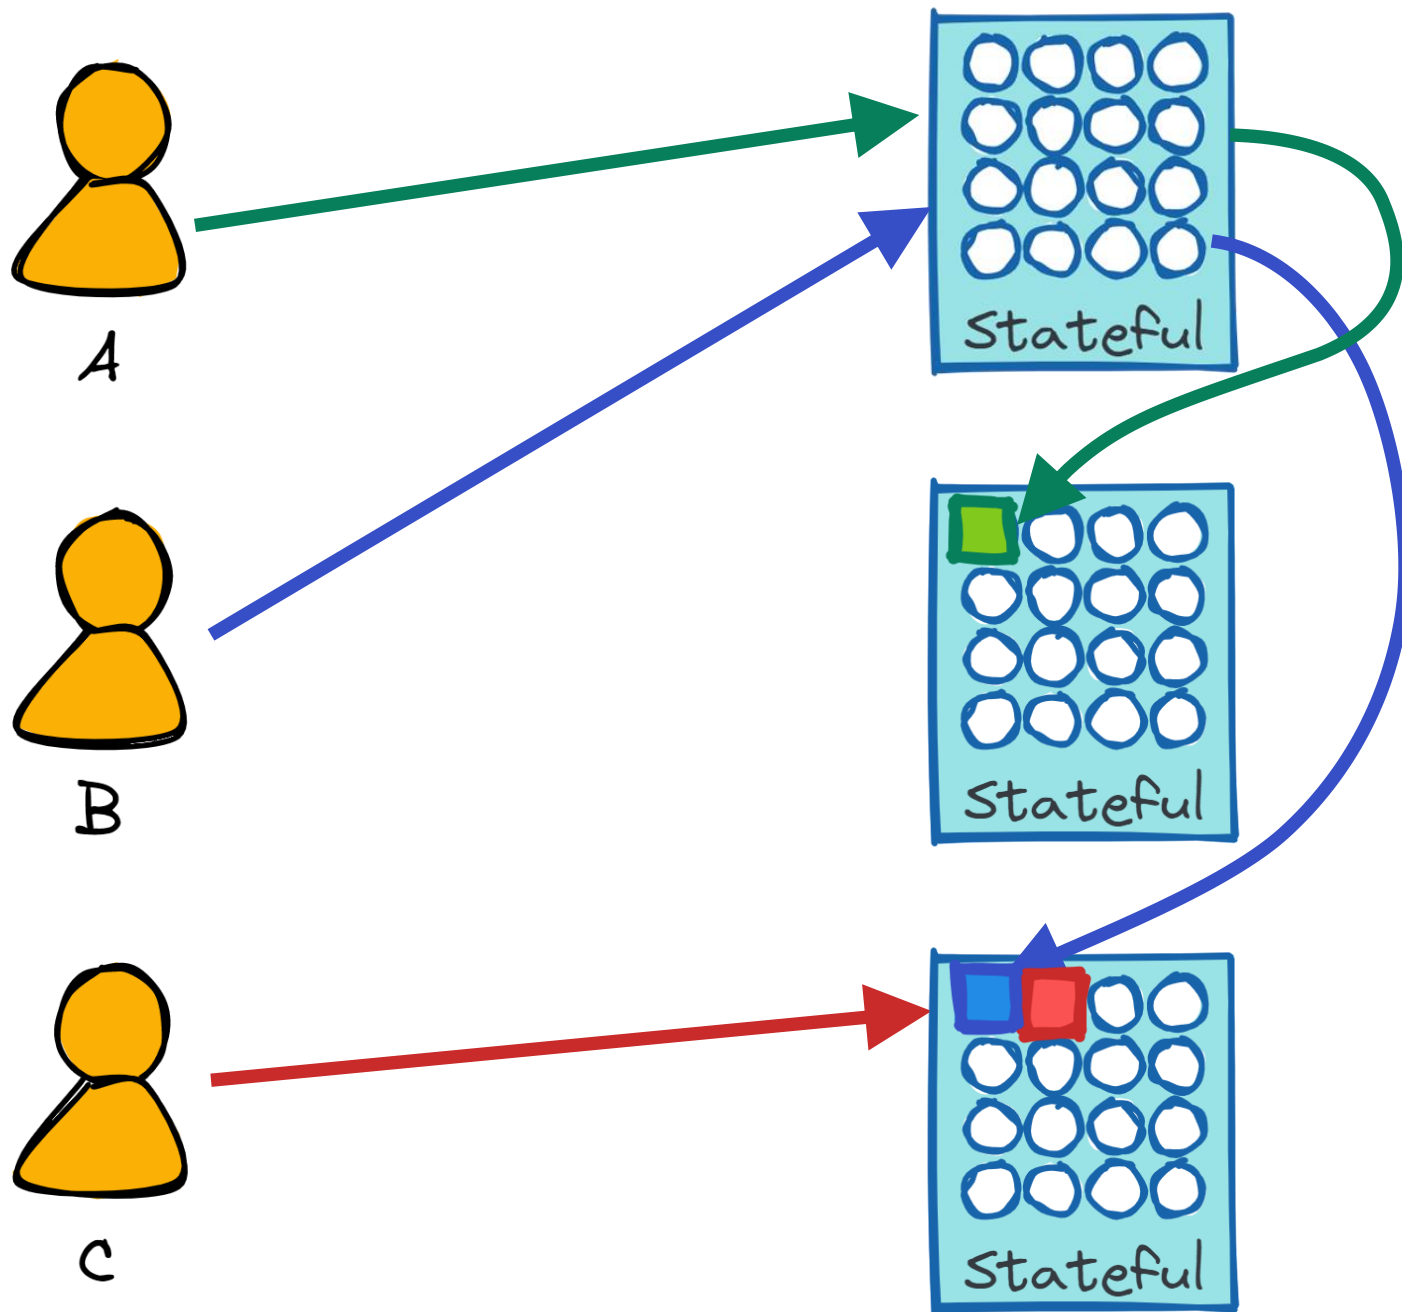
02

## The stateful solution

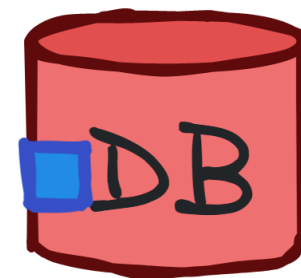
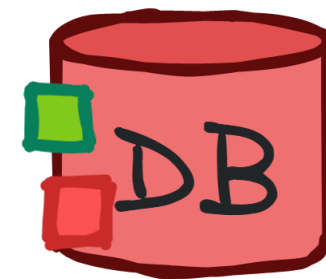




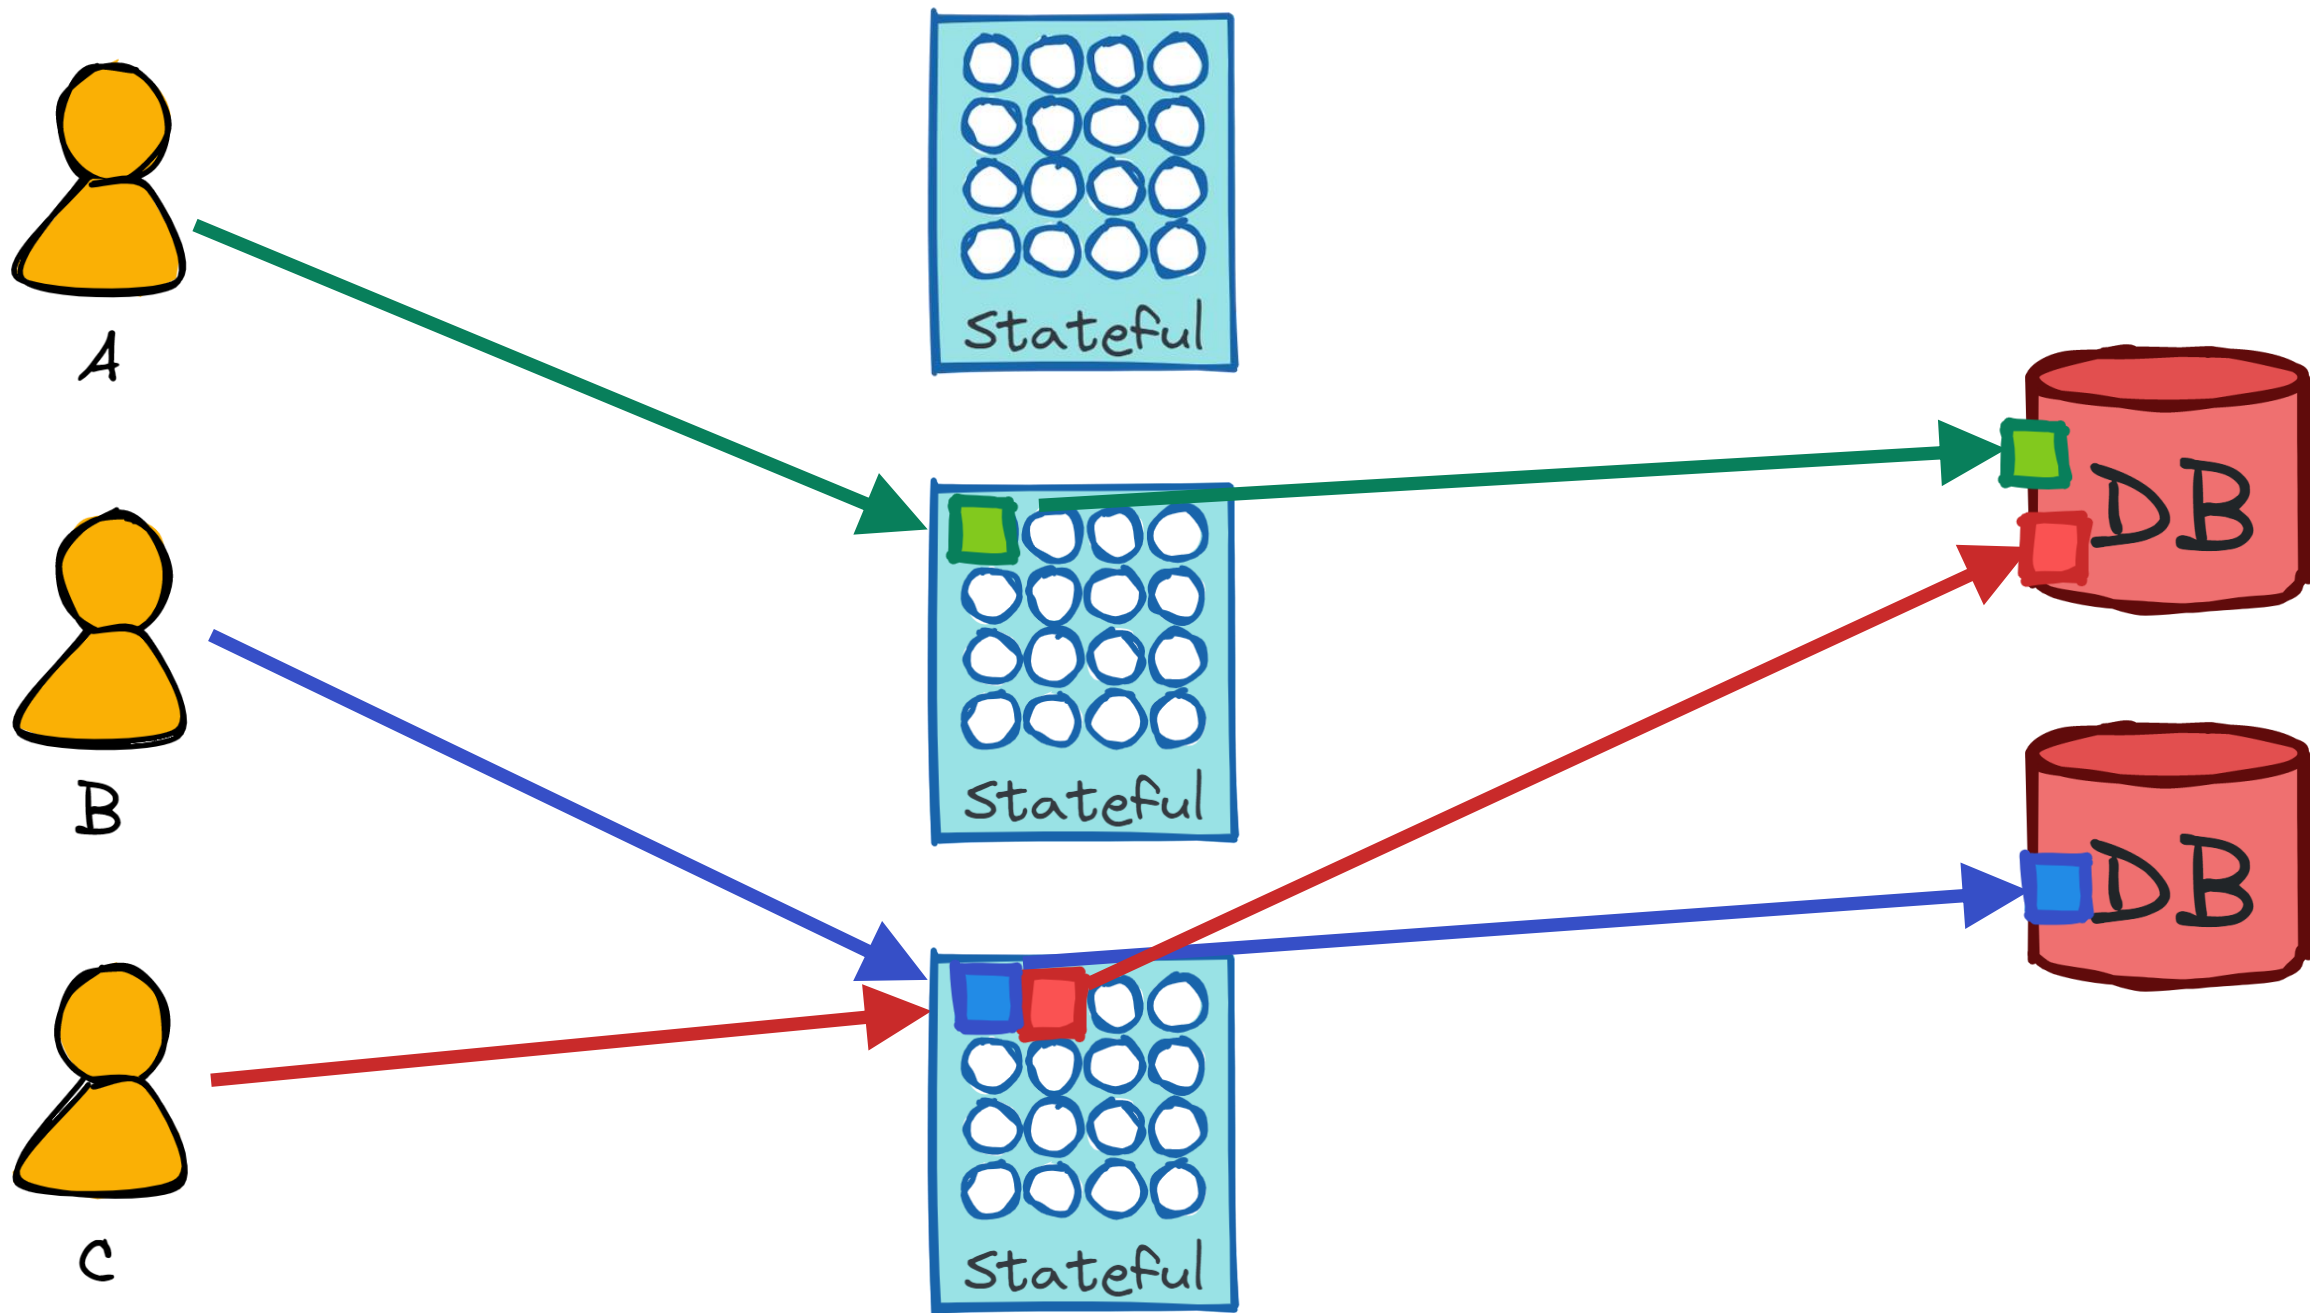


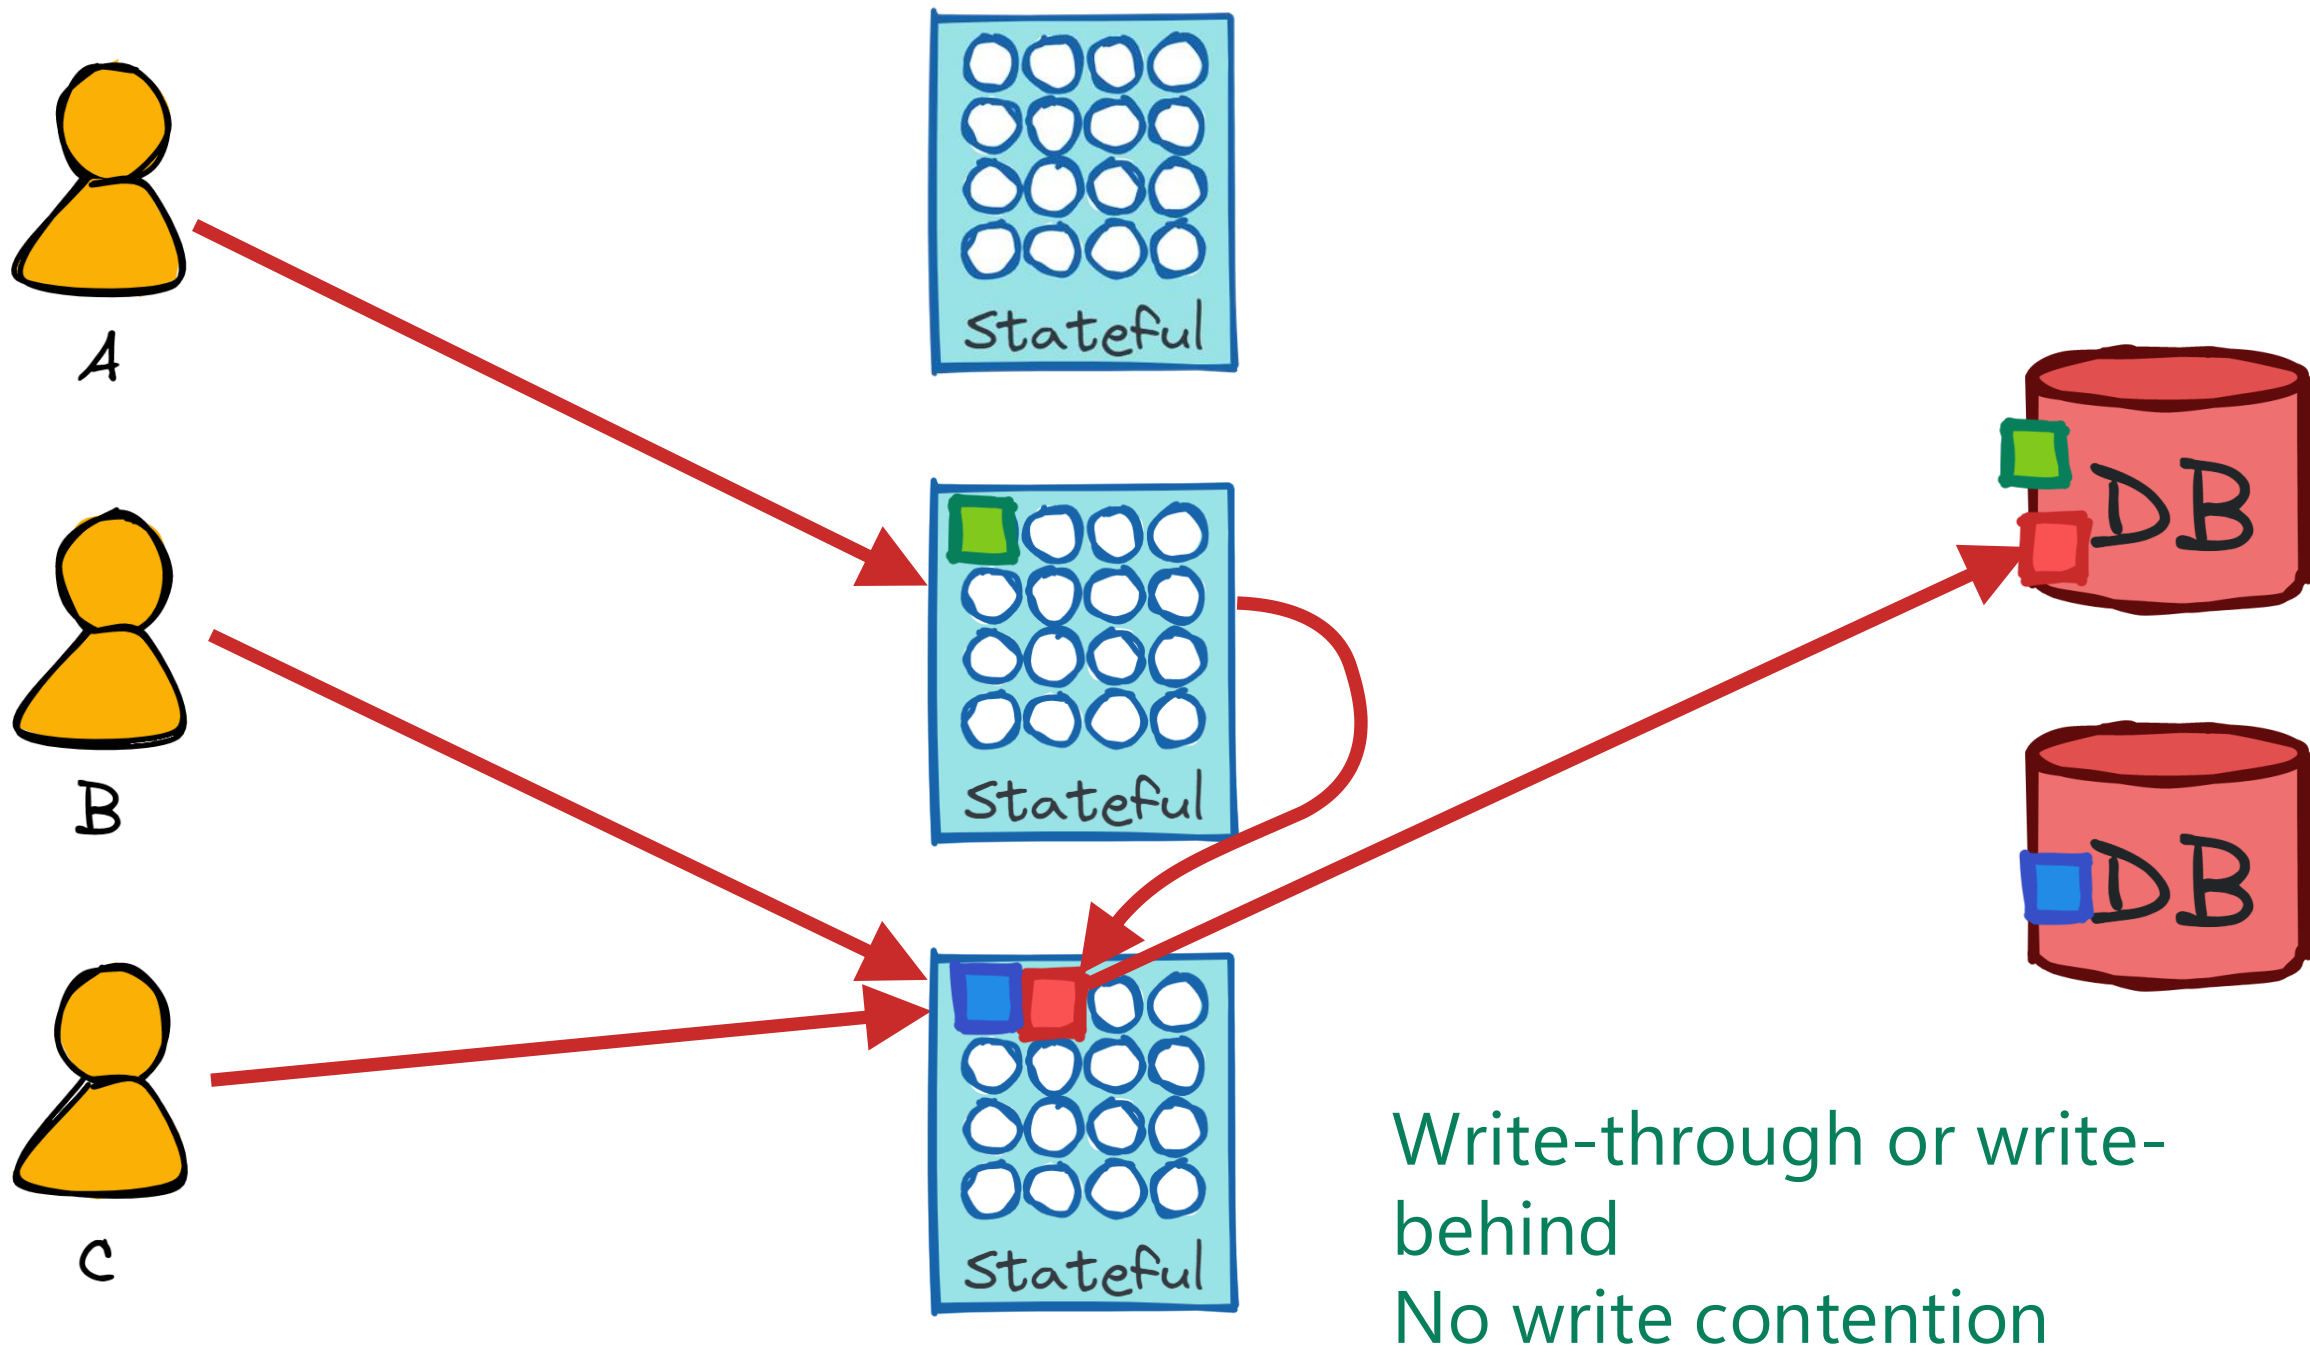


Fast, in-memory reads!









# Stateful services save cost, scale better

- Reduce database load
  - Hot/warm data is in memory, cold data is in database
- Eliminate need for separate cache servers
- Eliminate write contention
- Eliminate cache coherence issues
- Scales elastically
  - Load is balanced across your servers
  - More servers = more capacity
  - No partition keys

The background features several abstract purple geometric shapes. In the upper right, there are two large, curved, ribbon-like shapes that appear to be part of a larger structure. In the center, there is a solid purple sphere. In the lower left, there is a thin purple oval ring. The overall aesthetic is clean and modern with a monochromatic purple color scheme.

03

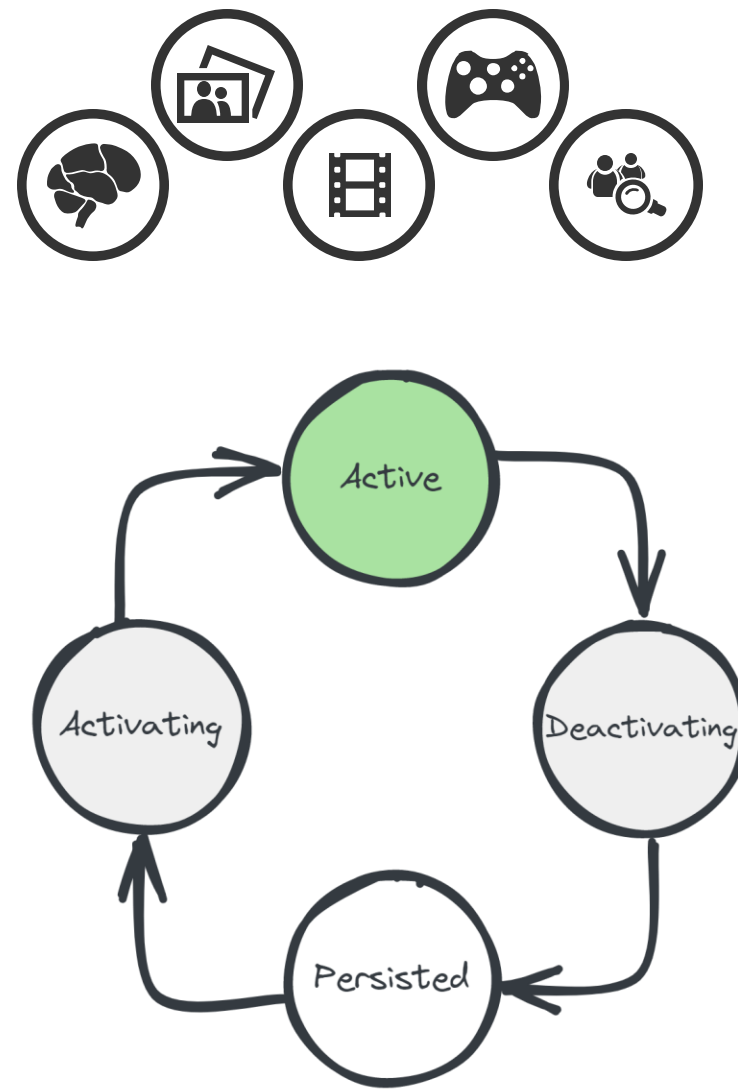
**Microsoft Orleans**

# Microsoft Orleans

- .NET library for (stateful) cloud applications
- Core primitive: grain (next slide)
- Runtime handles lifecycle, fault tolerance, load balancing, messaging, serialization, etc
- OSS, maintained by .NET team alongside Aspire, gRPC, ASP.NET
- Adopted by many Microsoft product groups & external companies, underpinning production services for 10+ yrs
- Continually evolving & improving (later in this talk)

# Grains

- Objects
  - Stateful, optionally persisted
  - Single-threaded by default
- Identified by **type** + **key**
  - Eg: **user/fred88**
- Loaded on-demand
  - Runtime managed lifecycle like virtual memory
- Communicate by RPC
  - Low-ceremony C# interface calls



```
public interface IUserGrain : IGrainWithStringKey
{
    Task SendMessage(IUserGrain from, string msg);

    Task<List<Message>> GetMessages();
}
```

```
public class UserGrain : Grain, IUserGrain
{
    IPersistentState<List<Message>> _msgs;

    public UserGrain([PersistentState("msgs")] IPersistentState<List<Message>> state)
        => _msgs = state;

    public Task<List<Message>> GetMessages() => Task.FromResult(_msgs.State);

    public async Task SendMessage(IUserGrain sender, string body)
    {
        _msgs.State.Add(new Message { Sender = sender, Body = body });
        await _msgs.WriteStateAsync();
    }
}
```

State is loaded  
by the runtime.

Methods are single-threaded.  
No locks necessary.



```
// Get grain references. Synchronous & cheap, like a URL
IUserGrain me = client.GetGrain<IUserGrain>("reuben");
IUserGrain friend = client.GetGrain<IUserGrain>("aditya");

// Call them
await friend.SendMessage(me, "hey :)");
var messages = await me.GetMessages();
```

# Other Features

## Persistence

Plugins for Azure Storage, Cosmos DB, SQL, DynamoDB, Redis, MongoDB, ...

## Pub/Sub

Decouple producers & consumers. Backed by Azure Event Hubs, Amazon SQS, GCP, in-memory

## Scheduling

Ephemeral timers, persistent reminders for scheduling future event

## Push notifications

Grains can send events (RPC calls) directly to clients. Eg, push notifications, streams

## Deploy anywhere

Azure Container Apps, Azure App Service, Kubernetes, AWS, GCP, bare metal/VMs, etc

## Custom placement & directory

Customize how the runtime decides where new grains are instantiated & how it finds them. “Move computation to data”

## Call filters

Wrap incoming & outgoing requests & responses for logging, tracing, error handling

## .NET integration

Aspire, Observability, Dependency Injection, Hosting, Configuration

The background features several abstract purple geometric shapes. In the top right, there are two overlapping curved bands. In the center, there is a solid purple sphere. In the bottom left, there is a thin purple oval ring. All shapes have a slight gradient and soft shadows on the white background.

04

**What's new since 8.0**

# What's new since 8.0

## Resource-aware placement

Places *new* grain activations on the least loaded hosts, with a locality bias

## Locality-aware repartitioning

Minimizes network traffic by migrating *related* grains into the same hosts

## Rebalancing

Rebalances already-active grains after scale-out, rolling upgrade, or changes in workload patterns

## Strong-consistency directory

Replaces the in-built grain directory with a strong-consistency, resilient directory

## Aspire integration

Streamlined inner-loop and “zero to cloud” experience using .NET Aspire

## Much more

Reliability & performance improvements, improved grain timers, support for MessagePack, Cassandra

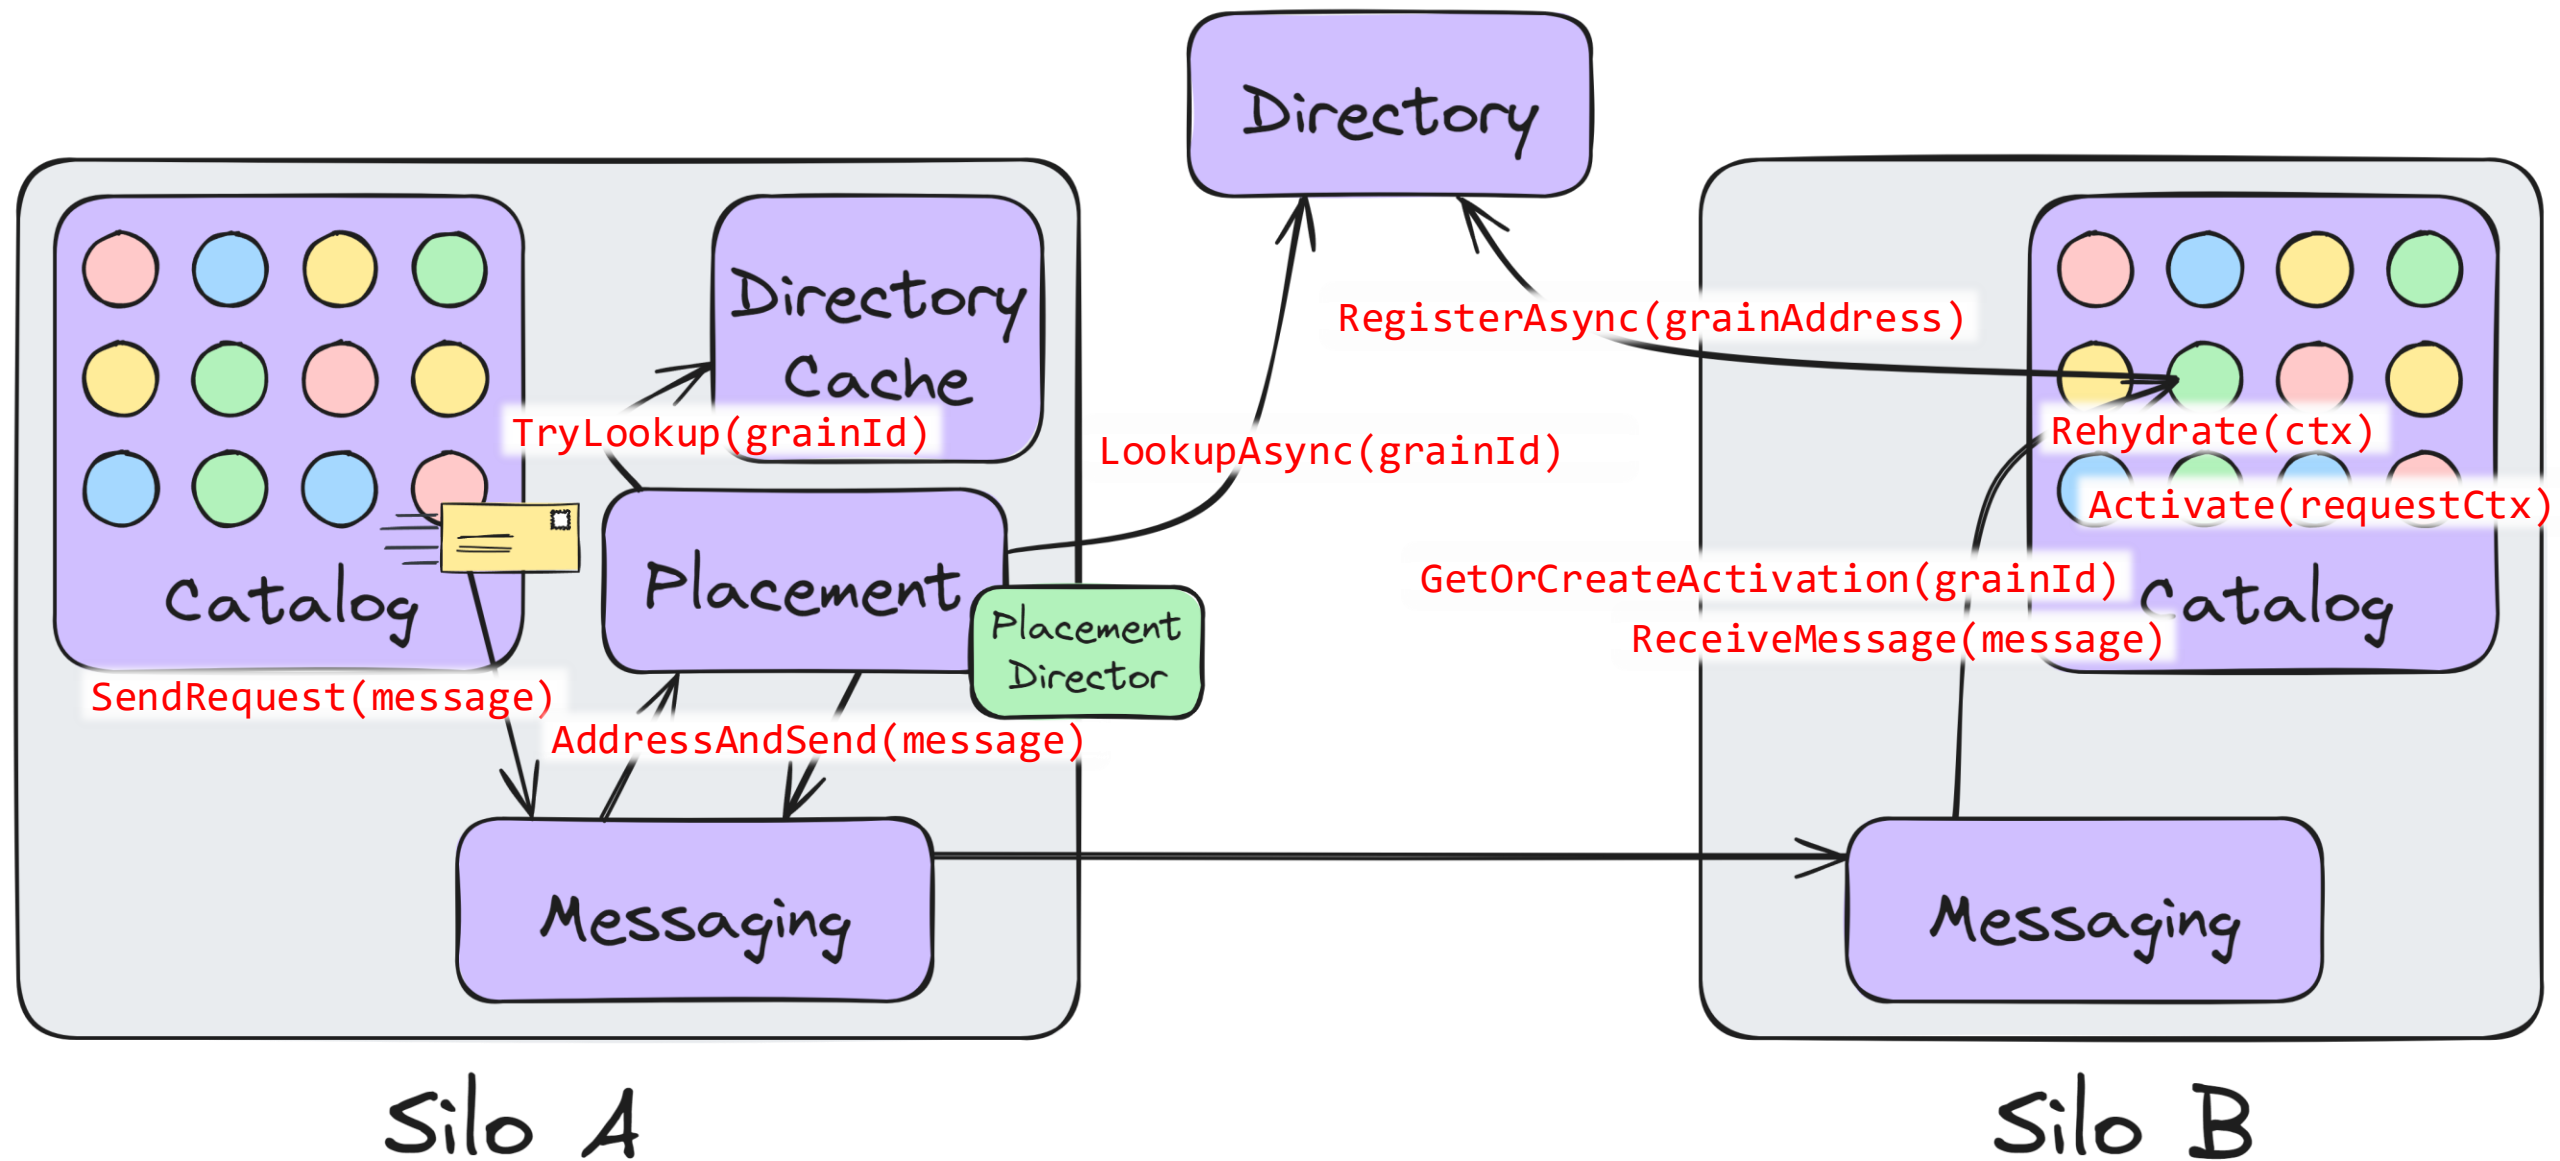
05

## Placement & Load Balancing

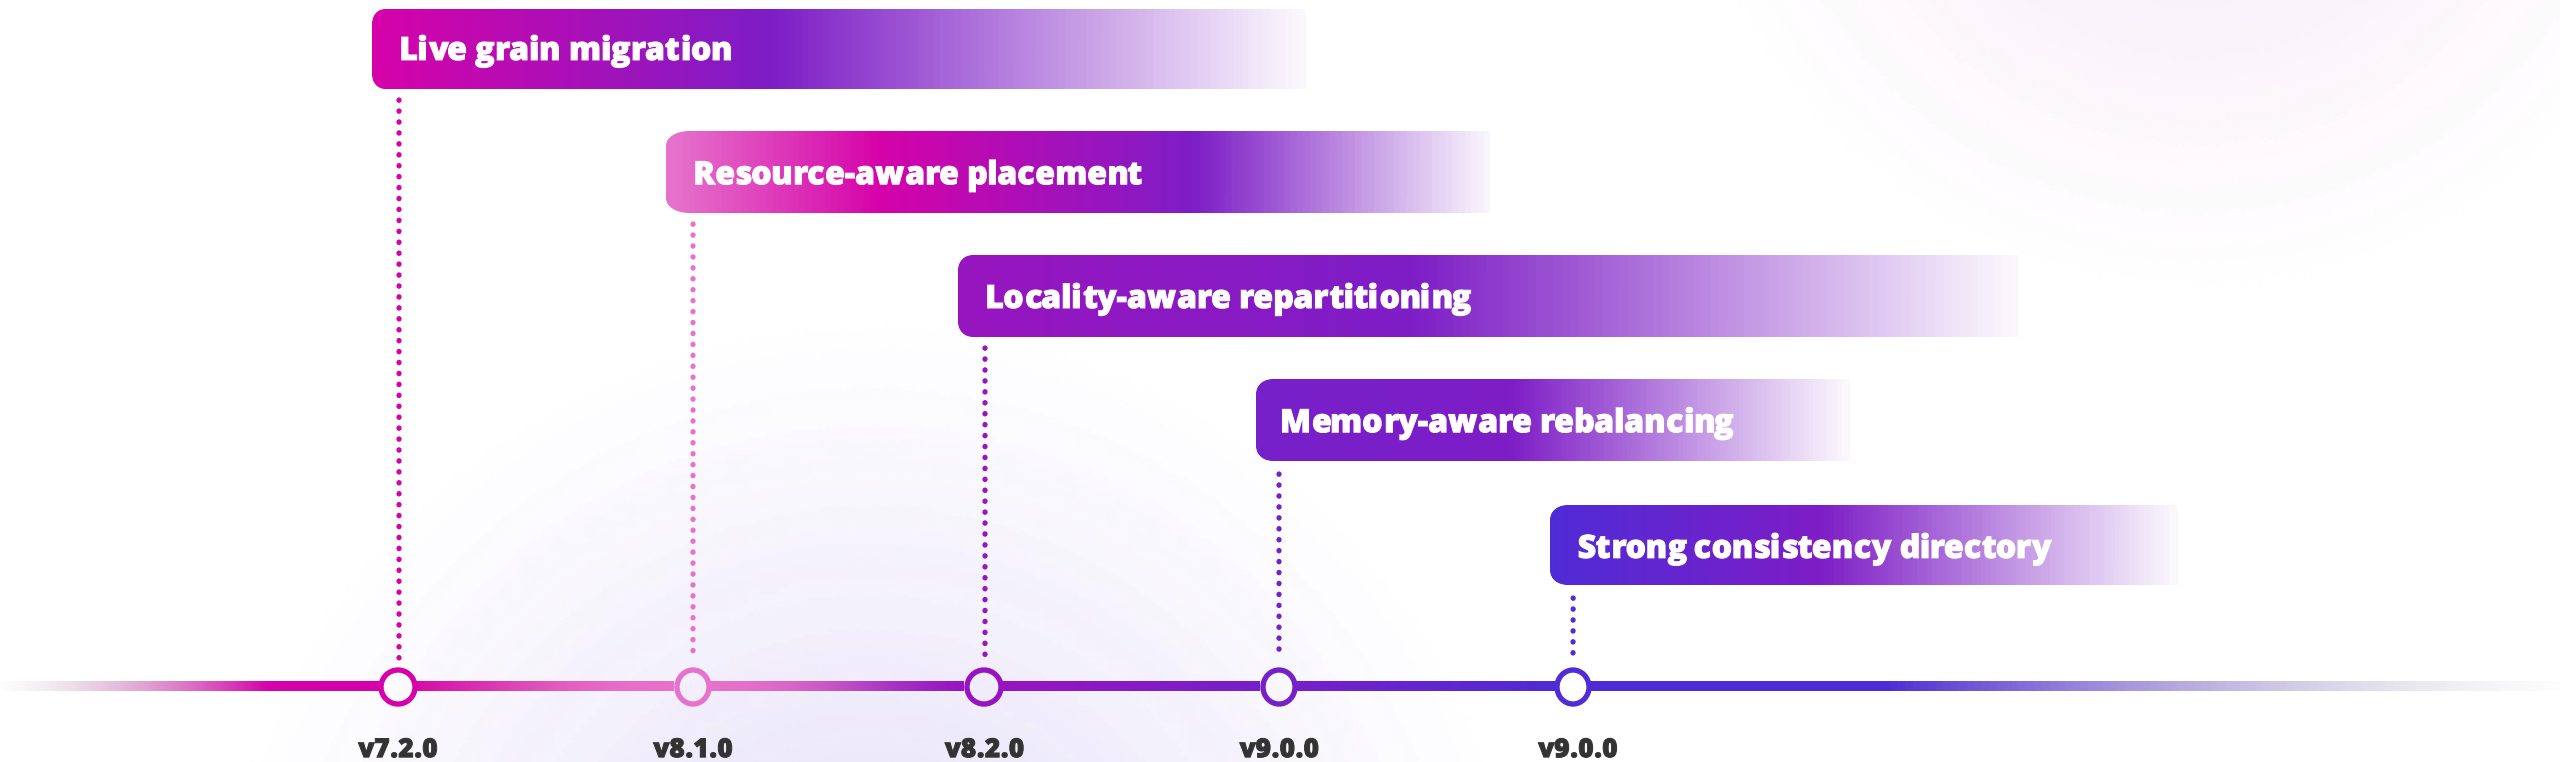


# Message routing overview

- Ask local directory **cache** which silo hosts the target grain
- If not found, ask **directory**, update cache
- If not found, run **placement** to pick a compatible silo
- Route message to that silo
- Silo **activates** target grain and **registers** it in the directory

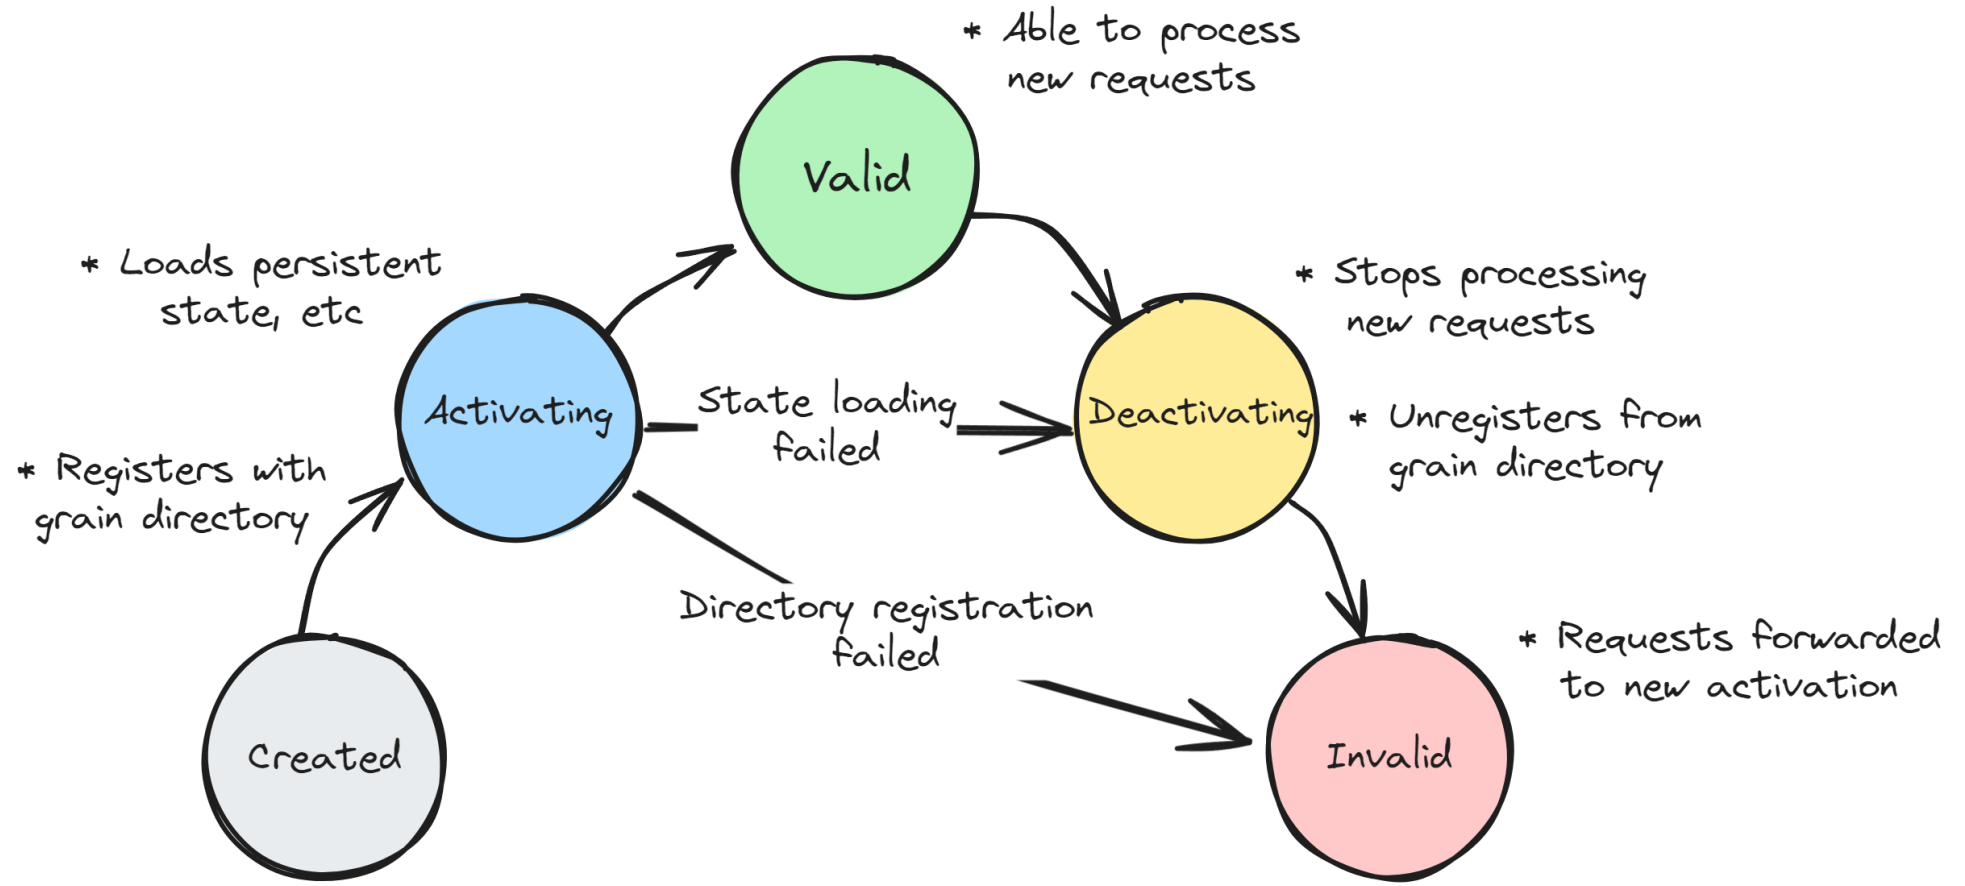


# Placement & load balancing enhancements

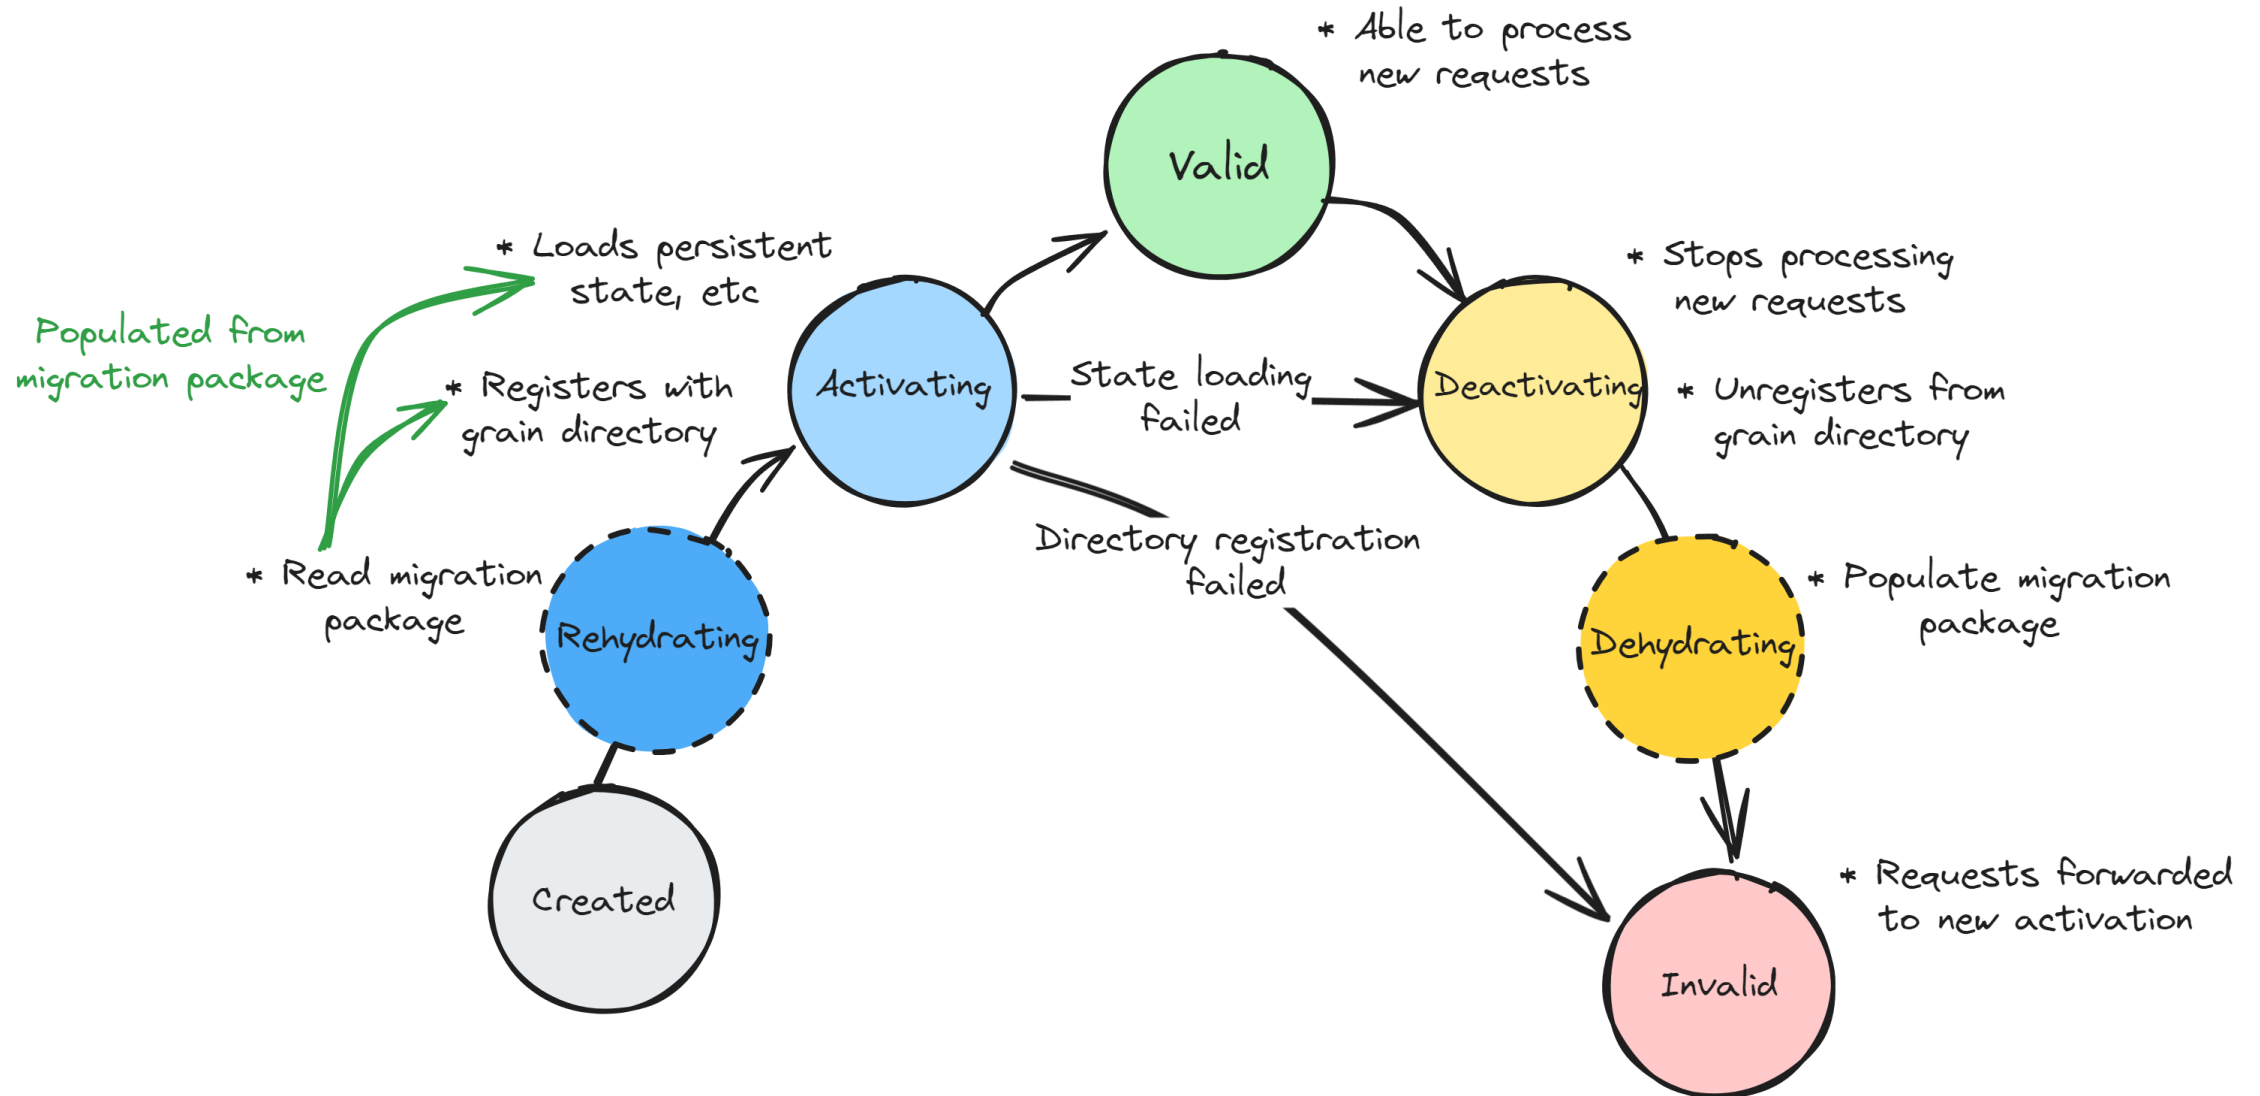




# Grain activation lifecycle

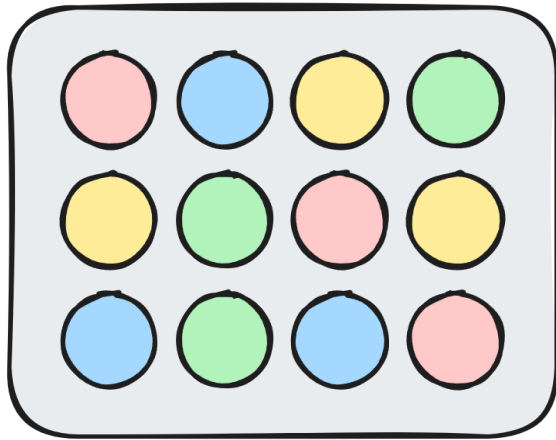


# ... with live grain migration

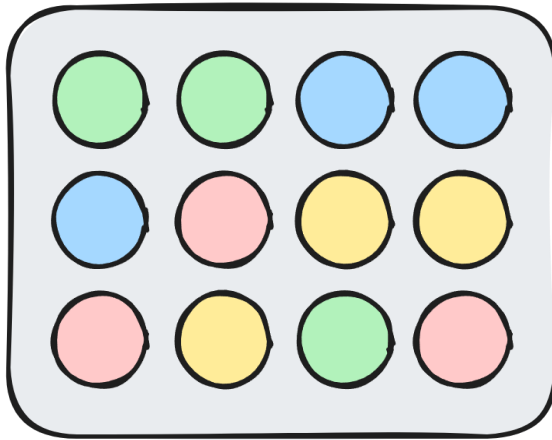


# Resource-aware placement

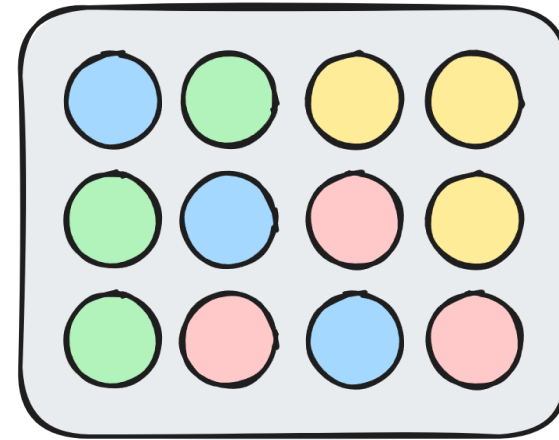
- Places activations on the least-loaded host
- Considers CPU usage, free/total memory, activation count
- Locality bias places related activations together



Silo A



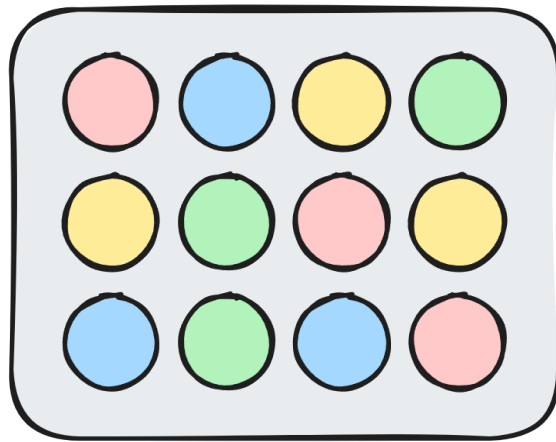
Silo B



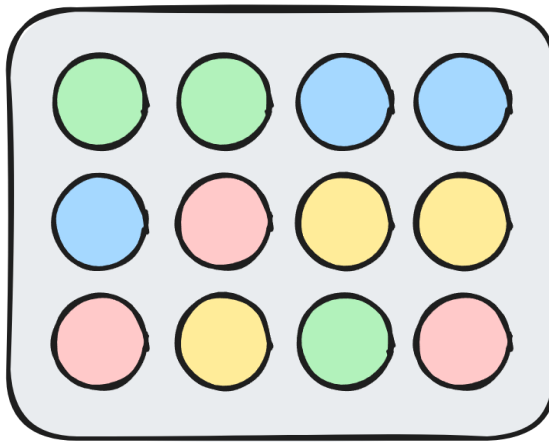
Silo C

# Locality-aware repartitioning

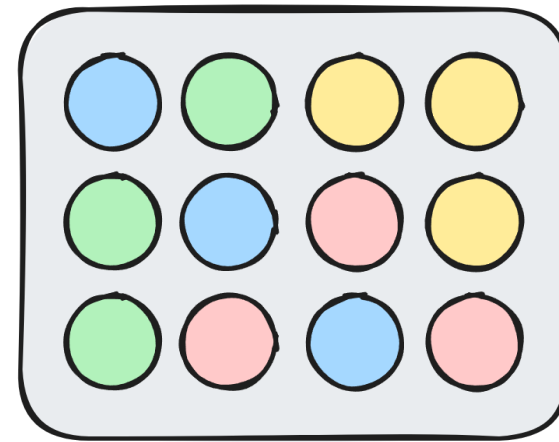
- Migrates *related* activations to the same host to reduce network calls, maintaining balance
- Improves throughput (**30-110%**), reduces overhead & latency



Silo A



Silo B



Silo C

Host 1 <sup>[1]</sup><sub>0</sub>

<sup>[4]</sup><sub>0</sub> Host 4

Host 2 <sup>[2]</sup><sub>0</sub>

<sup>[3]</sup><sub>0</sub> Host 3

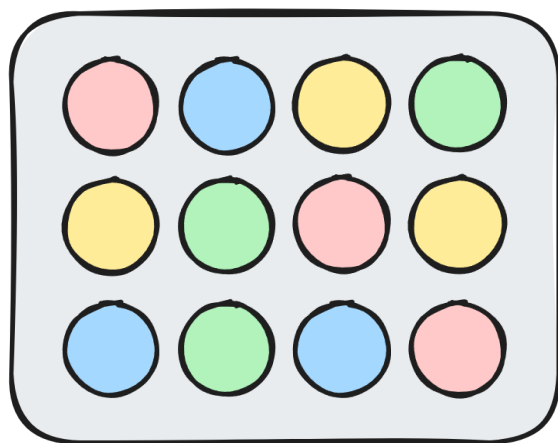
<sup>[0]</sup><sub>0</sub>  
Host 0

Adding load

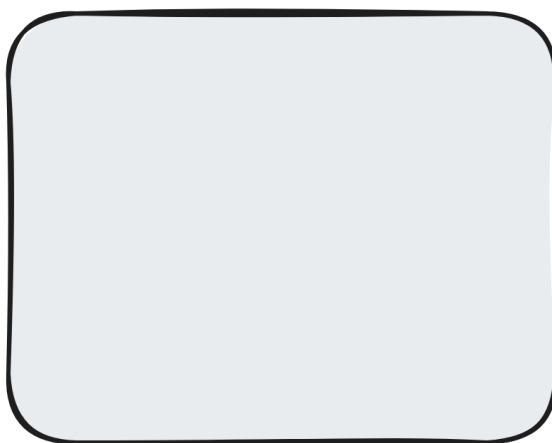
Red line = network call

# Memory-aware rebalancing

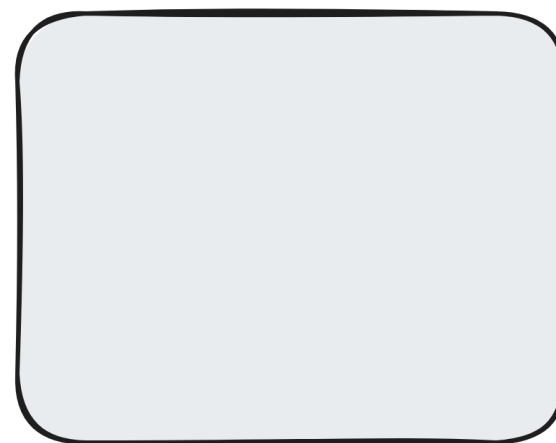
- Rebalance load as capacity increases
- In-memory state is transferred – no hit to database




Silo A

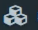



Silo B




Silo C


 Overview

 Grains

 Grain Details

 Silos

 Reminders

 Log Stream

Silos



ACTIVE SILOS

1



Silo Health

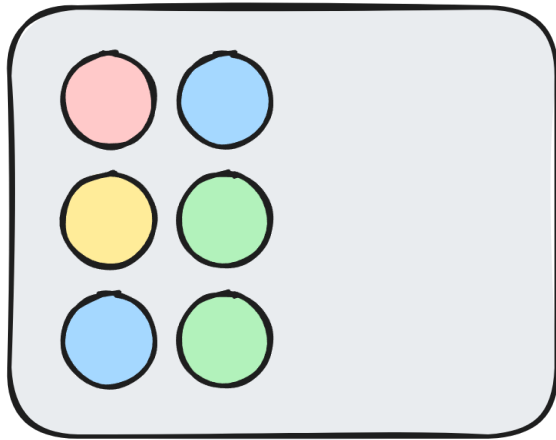
10.244.8.235:11111@88480289	Active	Silo_cb215	chaos-silo-859458c798-zlv6j	up for 2 minutes, 41 seconds	99736 activations
-----------------------------	--------	------------	-----------------------------	------------------------------	-------------------

Silo Map

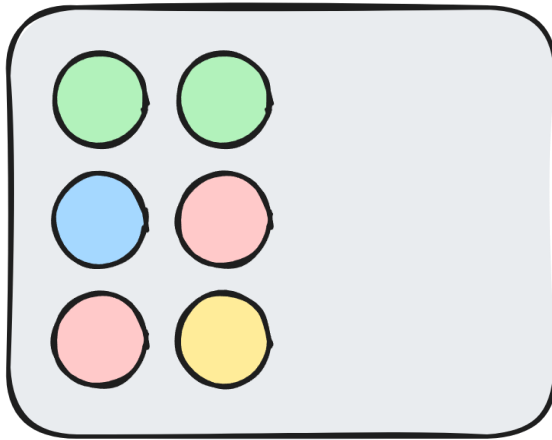
Update Zone 0		Fault Zone 0	
		10.244.8.235:11111@88480289 Active	

# Migrate on shutdown

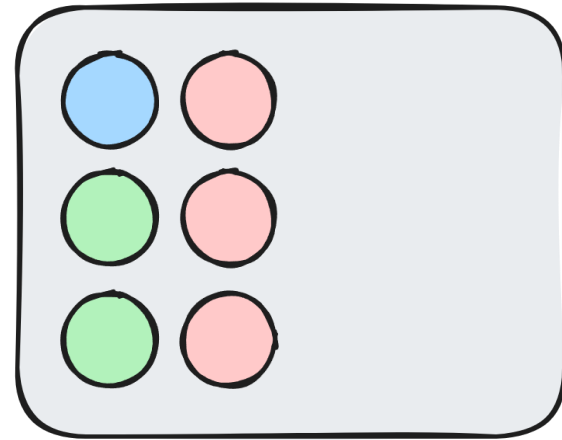
- Move valuable grains to surviving hosts during shutdown by calling `MigrateOnIdle()` from `OnDeactivateAsync(...)`



Silo A



Silo B



Silo C

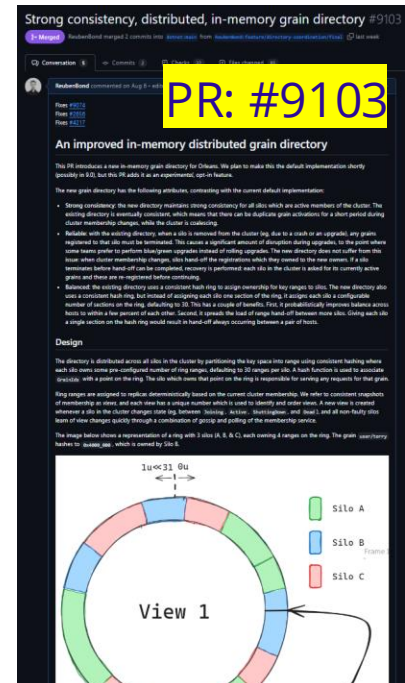


```
public sealed class MyRedGrain : Grain, IMyRedGrain
{
    public override Task OnDeactivateAsync(DeactivationReason reason, CancellationToken cancellationToken)
    {
        if (reason.ReasonCode == DeactivationReasonCode.ShuttingDown)
        {
            MigrateOnIdle();
        }

        return base.OnDeactivateAsync(reason, cancellationToken);
    }
}
```

# New in-cluster grain directory

- Original directory:
  - Eventual consistency, temporarily allows duplicate activations
  - Distribution prone to hot spots
  - Loses registrations during cluster churn, magnifying impact
- New directory:
  - Strong consistency
  - Well balanced
  - Reliable – no registration/activation loss
- Result
  - Smooth rolling upgrades, scale out/in



# Aspire integration

- Use Aspire to define infrastructure
- Aspire generates configuration for Orleans
- Orleans consumes configuration for various providers
  - Storage
  - Clustering
  - Reminders
  - Streams
- Aspire uses infra emulators/containers for local dev/test
  - Optionally provisions & configures real Azure infra
- Deploy to Azure Container Apps:
  - `azd init && azd up`

## AppHost project

```
var builder = DistributedApplication.CreateBuilder(args);

var redis = builder.AddRedis("orleans-redis");

var orleans = builder.AddOrleans("cluster")
    .WithClustering(redis);

builder.AddProject<Projects.DashboardToy_Frontend>("frontend")
    .WithReference(orleans)
    .WaitFor(redis)
    .WithReplicas(5);

builder.Build().Run();
```

## Frontend project

```
var builder = WebApplication.CreateBuilder(args);
builder.AddKeyedRedisClient("orleans-redis");
builder.UseOrleans();
var app = builder.Build();
```

# Special thanks to Ledjon Behluli

## Contributions & collaborations:

- Resource-aware Placement
- Locality-aware Repartitioning
- Memory-aware Rebalancing
- ...and more!



<https://github.com/ledjon-behluli>

Thank you, Ledjon!



[github.com/dotnet/orleans](https://github.com/dotnet/orleans)



[aka.ms/orleans/discord](https://aka.ms/orleans/discord)

**Thank you**